

When AI-Drafted Employment Records Become Evidence

By Hekim Colpan and Phillip Wikes

Both authors contributed equally. Names appear in alphabetical order. Hekim Colpan contributes in a personal professional capacity; the views expressed are his own and do not represent the position of any employer or institution. Phillip Wikes developed the Justification Review Standard described below.

When the stated reason becomes the issue

An employment decision does not end when a manager makes it. A performance evaluation, warning, promotion decision or termination memorandum may later become evidence through which someone asks whether the stated reason was legitimate, consistently applied and supported by what was known at the time.

When AI helps draft that record, the risk is not simply that the prose contains an error. It is that the record can become more polished than the evidence beneath it.

We call that gap Decision Reconstruction Risk: the condition in which a consequential record can no longer show, on its own, why the decision was made. In employment matters this is consequential because discrimination disputes can turn on the reasons an employer gives and the evidence surrounding them.

Pretext starts with the record

[*McDonnell Douglas Corp. v. Green*](#) established a burden-shifting framework. Where it applies, an employer may articulate a legitimate, nondiscriminatory reason for an employment action and the plaintiff may seek to show that reason is pretextual.

The Supreme Court did not hold that documentation quality determines whether discrimination occurred. Documentation can matter when the stated reason is tested against contemporaneous evidence, prior records and the consistency of the employer's explanation.

Consider a file where the underlying history includes emails, attendance records, completed work and prior feedback, while the final document contains a polished narrative that does not clearly connect its conclusions to those materials. An employer in that position can articulate its reason. What it may not be able to do is show the reason was the one it actually applied.

The decision may well have been the right one. The record may still be unable to demonstrate why it was made.

The pattern may appear only across employees

A single record is reviewed on its own terms. A workforce of records is reviewed together. That is where AI-assisted drafting introduces a risk most organizations do not currently measure.

AI-assisted drafting can reproduce language across performance reviews, disciplinary records and promotion decisions. A phrase that appears neutral in one file may take on a different

significance when it recurs across a group. The mechanism is straightforward: a drafting tool may reproduce similar characterizations when prompted with prior records, while a reviewer approving one record at a time may have no vantage point from which to notice the pattern.

Examples include "cultural fit," "executive presence," "not adaptable," "communication style" and "struggles with change." None establishes discrimination by itself. The issue is what happens when subjective language is repeatedly used to describe employees who share a protected characteristic and the organization cannot identify the evidence behind those descriptions.

Side-by-side review asks two questions: whether the same subjective standards are being applied across employees and whether the organization can identify the evidence supporting them. Disparate treatment and disparate impact remain distinct theories with different elements and proof structures. Recurring language does not establish either one by itself, but it can become relevant evidence when combined with surrounding facts, employment outcomes and decision history.

Banning particular phrases achieves little. What helps is requiring subjective conclusions to be connected to identifiable evidence before the record becomes final, and reading across records rather than only at each one.

What a defensible employment record should show

A consequential employment record should allow an independent reviewer to answer four questions:

- What happened?
- What evidence supports the characterization?
- Why did those facts matter?
- Was the same reasoning applied consistently?

"Not a strong cultural fit" is a conclusion. A record identifying specific conduct, dates, expectations and prior feedback gives a reviewer something that can be examined.

Before: "Attendance issues affecting the team."

After: "Missed nine scheduled shifts between January and March. Attendance expectations were discussed in documented feedback on two occasions."

Before: "Lacks professionalism."

After: "Missed client deliverable deadlines on March 3, March 10 and March 17. Feedback was provided after each occurrence."

The second versions make the basis of the judgment visible and give a later reviewer evidence against which the stated reason can be tested.

A pre-finalization control can close the gap

The control does not require a new platform or a wholesale redesign. It can be a structured review before a consequential employment record becomes final.

At minimum, the review should require the organization to:

1. Identify the human author and any AI tools used in drafting.
2. Preserve the source materials supporting material conclusions.
3. Link conclusions to verifiable source evidence rather than AI-generated assertions alone.
4. Document who reviewed the record, when and what substantive changes were made.
5. Define what drafting-layer material is preserved for consequential records and under which legal-hold triggers.
6. Review records across employees for recurring subjective language or inconsistent standards.
7. Confirm that the final explanation is consistent with documented history that existed before the decision.

The organizing principle is preservation rather than retention: enough evidence to reconstruct and defend the record when its author is no longer available to explain it.

A note for organizations operating in Europe

In Europe, the same employment-record problem sits at the intersection of data protection and AI governance. The [GDPR](#)'s accountability principle requires controllers to comply with the data-protection principles and to be able to demonstrate that compliance. It does not require every prompt or draft to be retained. But where personal data support an AI-assisted evaluation, warning, promotion or termination record, gaps in provenance, factual verification or human review can undermine the organization's ability to demonstrate compliance with the applicable data-protection principles.

[Article 22](#) adds a distinct safeguard where a decision is based solely on automated processing and produces legal or similarly significant effects. An employment record does not become an Article 22 case merely because AI helped draft it. That distinction makes meaningful human review more important, not less: formal approval is weak evidence of human judgment if the reviewer cannot identify the underlying facts, verify material conclusions or explain why the AI-assisted characterization was accepted.

Employment is also an area the [EU AI Act](#) treats expressly as sensitive. Annex III covers certain AI systems intended for recruitment and selection, decisions affecting promotion or termination, task allocation based on individual behavior or personal characteristics, and worker monitoring or evaluation. Whether a system falls within the high-risk regime depends on its intended purpose, not simply on AI appearing in the drafting process. Under [Regulation \(EU\) 2026/1744](#), the relevant high-risk requirements for Annex III systems are deferred until 2 December 2027. Many employment-drafting workflows may therefore fall outside that regime. **The governance question remains: when AI-assisted language becomes part of a permanent employment record, can the organization still show what evidence supported the characterization, what AI contributed, what a human verified and whether the same standard was applied consistently across employees?**

JRS as an operational example

One way to operationalize this is a structured review applied before a consequential record is finalized. The Justification Review Standard, developed by one of the authors, applies within

existing HR, compliance, investigations, audit and legal review processes. It asks whether an AI-assisted record can withstand independent review before it is finalized. It is undergoing structured validation using blinded reviewers, a predefined reference corpus and prespecified evaluation criteria, and results will be reported separately.

The instrument matters less than the discipline: any review that forces those questions before a record is finalized addresses the same failure.

The employment record is part of the decision

For compliance and HR leaders, the practical question is not whether AI should write employment records but whether the organization has a control at the point where AI-assisted language becomes part of the permanent record.

A defensible record lets someone who was not present reconstruct the reasoning and test whether the stated explanation matches the documented history.

When employment records are reviewed side by side, the question may shift from an individual employee's wording to whether the organization's records reveal standards that were subjective, inconsistently applied or difficult to defend.

We use the shorthand "right to know why" for that governance principle. It is not a legal doctrine and not a claim of any new entitlement. The tool doing the drafting has changed. The standard the record has to meet has not.

About the authors

Hekim Colpan is an AI Governance and Compliance Manager, Data Protection Manager and ISO/IEC 42001 auditor based in Germany. His work focuses on the operational implementation of AI management systems, the EU AI Act, GDPR accountability, DORA and governance controls for AI-assisted systems.

Phillip Wikes is an AI Governance and Cognitive Risk Advisor focused on documentation integrity, evidentiary traceability and record-level controls in AI-assisted environments. He served as a Lead Civil Rights Officer at the Maryland Commission on Civil Rights and developed the Justification Review Standard and Decision Reconstruction Risk. He holds an M.S. in Negotiation and Conflict Management.