

FINAL SUBMISSION READINESS AUDIT

Detection Article v3

Date. August 18, 2026

Manuscript version. Detection_Article_Submission_Final_v3_2026-08-18

Report version. Report version 1.0, 18 August 2026

Mode. Read only. No manuscript, dataset, script or pre-registration was modified.

Final status. AMBER

Recommendation. ONE SURGICAL CORRECTION THEN SUBMIT

Contents

EXECUTIVE FINDING

PART 1. REFERENCE-CLASSIFICATION PROVENANCE

PART 2. REFERENCE-KEY TIMELINE

PART 3. PARTICIPANT ARCHITECTURE

PART 4. PRIMARY STATISTICAL LINEAGE

PART 5. RELIABILITY STATISTICAL LINEAGE

PART 6. HISTORICAL VALUE CONTAMINATION

PART 7. JRS CLAIM AUDIT

PART 8. DRR CLAIM AUDIT

PART 9. REFERENCE-CLASSIFICATION CLAIM AUDIT

PART 10. REPRODUCIBILITY DISCLOSURE

PART 11. ETHICS / PARTICIPANT DISCLOSURE

PART 12. NOVELTY CLAIM

PART 13. CONSTRUCT-DEPENDENCE

PART 14. EDITORIAL OVER-DEFENSIVENESS

PART 15. UNRELATED TECHNICAL DEBT

PART 16. FINAL SUBMISSION GATE

FINAL RECOMMENDATION

EXECUTIVE FINDING

v3 is submittable after one surgical correction. There is no methodological defect requiring repair, and no reported statistic fails to trace to an authoritative source.

The v3 pass closed the one issue that was a genuine blocker: the manuscript no longer presents an automated reproducibility check as human validation. All fourteen instructed corrections hold, and 59 standing assertions plus 23 zero-drift checks pass against the current file.

One item warrants a correction before submission. The novelty claim at Section 2 is an unhedged universal negative about the literature: *"reconstructability of the individual record has not been operationalised as a measurable property with a stated instrument, a stated scale, and reported detection and agreement statistics."* No source can establish a universal negative, and reviewers in this field routinely require "to our knowledge" or equivalent. This is a one-clause fix, not a methodological problem.

Three items are AMBER and require your judgement rather than an edit. The most consequential is the claim that the author-side classification was fixed *"before any reviewer was recruited"*. It is not contradicted by anything, but it is also not verifiable from the repository: `pilot_progress` exposes no registration timestamp. You know whether it is true; the repository cannot show it either way.

The reproducibility position is honestly stated and is weaker than the paper's other evidence. The 72 reference classifications are supported by an aggregate result, not by retained per-pass outputs. v3 now says so. That is a real limitation, correctly disclosed, and not a submission blocker.

PART 1. REFERENCE-CLASSIFICATION PROVENANCE

1. Are the Section 4.4 instances the same as the Appendix A models?

NOT ESTABLISHED.

`research/Verified_Key.md` records that three large-language-model instances performed the pass and names no vendor, model or version. `Detection_Article...:220` names three vendors (Anthropic, OpenAI, Google) for the Appendix A nightly cross-vendor runs. **No source connects the two procedures.** The inference is available and is not drawn.

2. If SAME

Not applicable. For completeness, the retained configuration fields for the reference pass are:

Field	Retained
Model name	No
Vendor	No
Version	No
Execution date	No (the file was committed 2026-07-06; the pass itself is undated)
Prompt	Partially — the operational rule given to the raters is <code>research/AnswerKey_Verification_Packet.md</code> ; no system prompt or wrapper is recorded
Temperature or sampling configuration	No
Per-pass outputs	No

3. If DIFFERENT

Not established as different either. The two procedures are distinct in purpose and are distinguished in the manuscript:

	Reference classification, Section 4.4	Machine consistency, Appendix A
Purpose	reproduce the author-side key	measure cross-vendor agreement
Instrument	binary operational rule	five-condition JRS instrument
Corpus	24 records	15 records
Vendors named	none	three, named
Reported as	methodological check	supporting result

The manuscript distinguishes them adequately. They sit in different sections, describe different instruments and different corpora, and neither cross-references the other in a way that could merge them.

4. If NOT ESTABLISHED

The manuscript's current wording remains accurate. Section 4.4 says "three separate large-language-model instances" and asserts nothing about identity. No model name was added for symmetry with Appendix A.

5. Are the 72 classifications supported by retained outputs?

NO. The arithmetic is supported; the outputs are not retained.

Element	Status	Source
3 instances	stated	Verified_Key.md, Method
24 records	tabulated	Verified_Key.md, key table, R01–R24
$3 \times 24 = 72$	arithmetic, verified	—
72 individual retained classifications	not retained	no per-pass artifact exists

What exists is the aggregate: *"All three raters assigned identical labels on all 24 records."* v3 discloses this limitation in the Disclosure paragraph.

6. 72/72 agreement

VERIFIED as reported, from the aggregate. Verified_Key.md, Result: 100 percent inter-rater agreement, and 24/24 against the intended key. Unanimity across three passes on 24 records entails 72 agreeing classifications. Not independently re-derivable from retained per-pass data.

7. Adjudication

VERIFIED: none occurred. No divergence, so the pre-specified rule at OSF_PreRegistration.md:28 was never triggered.

8. Pre-registered 2, executed 3

VERIFIED. OSF_PreRegistration.md:27–28, DRR_Detection_Validation_Protocol.md:30–33, AnswerKey_Verification_Packet.md:16 and Intended_Key_authorside.md all specify two with conditional adjudication. Verified_Key.md records three, each scoring all 24.

9. Unreported methodological consequence of the 3-pass procedure

One, and it is minor. The pre-registered design was one full pass plus a conditional adjudicator. The executed design was three full independent passes, which permits a majority rule the pre-registration did not contemplate. Because agreement was unanimous, no rule of any kind was needed and the two designs coincide on this data. Had there been divergence, the executed design would have required a resolution rule that was never specified. **No consequence for the reported result.**

10. Is the current disclosure sufficient?

YES. v3 states the count, the automated nature, the absence of expert status, the denominator, the packet's disclosure of the task, the absence of adjudication, the deviation from two passes to three, that the check is not human validation, that no human replication has been performed, and that implementation details were not retained in a form sufficient for independent reproduction.

PART 2. REFERENCE-KEY TIMELINE

	Event	Date	Source
A	Author-side key created	on or before 2026-07-06 ; creation itself undated	Intended_Key_authorside.md; git log db7a34c
B	Fixed and committed	2026-07-06	git log db7a34c, 48f3ead
C	Detection panel reading	began on or before 2026-06-28	live pilot_progress, earliest last_at 2026-06-28T08:33:02Z
D	Study 004 labelling	began 2026-06-11	live bench_labels, earliest created_at 2026-06-11T21:01:51Z

	Event	Date	Source
E	Primary analysis	2026-08-15	research/FULL_DATA_ANALYSIS_2026-08-15.txt; closed_aggregates_2026-08-15.json generated_at 2026-08-15T06:27:02Z
F	Pre-registration requirement	"The verified key ... is fixed before any accuracy analysis "	OSF_PreRegistration.md

The pre-registration requirement is met. It governs analysis, not recruitment or reading. Analysis ran on 2026-08-15, six weeks after the key was fixed. Reviewers never saw either key, so blinding is unaffected.

The "before any reviewer was recruited" claim

Detection_Article...:172: *"Before any reviewer was recruited, the first author recorded an intended classification..."*

PARTIALLY VERIFIED.

- **Not contradicted.** Nothing in the repository shows a recruitment predating the key.
- **Not verifiable.** `pilot_progress` exposes code, `last_at`, `name`, `reads_today`, `total_reads` and no registration timestamp. `ai_pilot_reads` is RLS-locked and returns zero rows to the public key. No recruitment or invitation log exists in the repository.
- **Circumstantially strained.** The key file was committed 2026-07-06; a participant's final activity is recorded 2026-06-28, so reading had begun at least eight days earlier. The claim can still be true — the intended key states it was fixed before verification, and a file can be authored well before it is committed — but the repository cannot demonstrate it.

Classification: EDITORIAL RISK, not a submission blocker. The manuscript does not rest any statistical claim on it, and the substantive protection against post-hoc fitting is the blind reproduction, not the recruitment sequence. You know from personal knowledge whether the claim is true; if it is, no change is needed. If you cannot date it confidently, the safe form is *"before verification began"*, which the intended-key file does establish.

PART 3. PARTICIPANT ARCHITECTURE

Population	Count	Nature	Verified from
Study 011 / Arm A	16 completers of 27 registered	independent human experts	live <code>pilot_progress</code> ; <code>Expert_Roster_All_Studies_2026-08-06.md</code> Study 011
Study 012 / Arm B	20 completers of 21 registered (B1 7, B2 13)	independent human experts, same professional standing	live <code>armb_progress</code> ; same roster, Study 012
Study 004 reliability	25 raters (8 E-, 17 R-)	human	live <code>bench_labels</code> ; <code>REVIEWER_ROSTER_COMPLETE.md</code> §004
Reference classification	3	automated LLM instances	<code>Verified_Key.md</code> :8

Check	Result	Evidence
Arm A expertise	VERIFIED	all 16 named with credentials
Arm B expertise	VERIFIED	20 completers; anonymous entries typed "JRS-naive expert professional"
Arm A / Arm B distinction is condition, not expertise	VERIFIED	<code>DRR_Detection_Validation_Protocol.md</code> :46
Arm A / Arm B disjoint	VERIFIED	no named person holds both a V-AI and an RR code
Study 004 kept separate	VERIFIED	protocol §4 names no E- or R- code
LLM separation from all human populations	VERIFIED	zero overlap; not human
61 human study participations	VERIFIED	16 + 25 + 20
58 unique humans	VERIFIED	61 – 3 dual-code holders
Three-person overlap	VERIFIED	E-09/V-AI-06, E-12/V-AI-07, E-13/V-AI-03, <code>build_expert_roster.py</code> :121

No manuscript statement conflicts with this architecture. The Acknowledgments are byte-identical to v2 and credit 61 participations held by 58 distinct people. The three automated raters appear nowhere in that count.

One residual, recorded rather than smoothed over. E-11 is one of the 25 credited raters and carries one label with no identity row, so it cannot be resolved to a named person. `count_participants.py` reaches 58 by a different composition that excludes E-11 and includes RR-108, a non-completer; the two exclusions cancel. Both routes print 58 while describing sets that are not identical. No reported figure depends on which route is taken.

PART 4. PRIMARY STATISTICAL LINEAGE

Statistic	Source data	Analysis file	Output file	Manuscript location	Verified
83.9% accuracy	ai_pilot_reads scored against research/Verified_Key.md	api/pstat-4c8e1b6a2d90.js (deleted at 120c11e, recovered from 120c11e^)	research/closed_aggregates_2026-08-15.json → detection_panel.accuracy.mean = 83.85	Abstract, §6.1	YES
95% CI 72.7–95.1	same	same, tcrit(15)=2.131, h = t · s / √n	same → ci95_low 72.66, ci95_high 95.05	Abstract, §6.1	YES
Sensitivity 87.0%	same	same	same → sensitivity.mean = 86.98	§6.4	YES
Specificity 80.7%	same	same	same → specificity.mean = 80.73	§6.4	YES
384 graded reads	same	same	same → judgments_analysed = 384	Abstract, §6.1, Acknowledgments	YES
16 experts	pilot_progress	research/check_completion.py; research/count_participants.py	live read 2026-08-18: 27 registered, 16 at ≥24 reads	throughout	YES
24 records	research/Verified_Key.md key table	—	R01–R24, 12 grounded / 12 ungrounded	throughout	YES

The confidence-interval method was verified arithmetically: $2.131 \times 21.02 / \sqrt{16} = 11.198$; $83.85 \pm 11.198 = 72.65 / 95.05$, matching the stored values to the rounding of the stored SD, and printing as 72.7 and 95.1. Corroborated independently by `scripts/verify_detection_accuracy.py`, whose `ci95_t()` carries the same t table.

One lineage weakness, already on the register. The producer `api/pstat-4c8e1b6a2d90.js` was deleted after use and exists only in git history. The output file names its own producer, and the arithmetic reproduces, so the lineage is traceable — but it is traceable through `git show`, not through a file a third party would find in the working tree.

PART 5. RELIABILITY STATISTICAL LINEAGE

Statistic	Source data	Analysis file	Output file	Manuscript location	Verified
Expert AC1 0.739	live bench_labels, mode = jrs, deduplicated	research/compute_ac1_ci.py ac1()	research/current_reliability_2026-08-18.json	§6.5 table	YES
Expert analytic CI 0.402–1.000	same	compute_ac1_ci.py analytic_ci()	same	§6.5 table	YES
Expert bootstrap CI 0.427–1.000	same	compute_ac1_ci.py bootstrap_ci(), B=20000, seed=20260727	same	§6.5 table	YES
Regular AC1 0.623	same	ac1()	same, 0.6228	§6.5 table	YES
Regular analytic CI 0.252–0.993	same	analytic_ci()	same	§6.5 table	YES
Regular bootstrap CI 0.285–0.894	same	bootstrap_ci()	same	§6.5 table	YES
25 / 22 / 3 participants	live bench_labels	research/recompute_current_ac1.py	same, totals and excluded_baseline	§6.5, Acknowledgments	YES
113 / 104 determinations	same	same	same, 113 submitted, 104 retained	§6.5, Appendix B	YES

Statistic	Source data	Analysis file	Output file	Manuscript location	Verified
15 / 10 records	same	same	same, 15 labelled, 10 estimable, 5 single-rater	\$6.5	YES

Is `compute_ac1_ci.py` responsible for reported statistics?

YES. It is the estimator behind every reliability value in the manuscript. `research/recompute_current_ac1.py` imports it unmodified and calls `ac1()`, `analytic_ci()` and `bootstrap_ci()`, reading that module's own `B` and `SEED` constants at run time.

Consequence of the crash

D. REPRODUCIBILITY ONLY.

`compute_ac1_ci.py` raises `KeyError: 'basis_identification'` at `per_condition_ac1()`, line 169, when run as `__main__` against its own committed dataset `research/construct_validity_data.csv`, which carries three columns where the function needs five.

Component	Reported in the manuscript	Affected by the crash
<code>ac1()</code>	yes, both AC1 values	no
<code>analytic_ci()</code>	yes, both analytic intervals	no
<code>bootstrap_ci()</code>	yes, both bootstrap intervals	no
<code>krippendorff_alpha()</code> , <code>fleiss_kappa()</code>	not reported numerically in v3	no
<code>per_condition_ac1()</code>	not reported anywhere	crashes

No reported value passes through the crashing function. The defect is that a replicator who runs the script as documented gets a traceback before seeing the trained-reviewer block, and must import the module to reach it, as this audit did. It affects the replication experience, not the results.

A second defect in the same file: `dedup_last()` is a no-op against the committed CSV, which was already deduplicated before commit. Consequence: **no committed file demonstrates the 113 → 104 reduction.** The reduction is verified against the live table instead, and `recompute_current_ac1.py` sorts by `created_at` before deduplicating so the "latest submission retained" rule is deterministic rather than iteration-order dependent.

PART 6. HISTORICAL VALUE CONTAMINATION

SUPERSEDED VALUES FOUND: 0.624, 0.253 to 0.994, 0.301 to 0.886, 63 labels, 108 submitted, 99 after keeping, plus the retired terminology trained reviewer, blind raters, blinded raters, same pool, those same experts, expert panel, All 61, 36 independent experts.

SUPERSEDED VALUES IN ACTIVE MANUSCRIPT: none. `scripts/verify_manuscript_figures.py` checks 29 superseded values against v3 and reports the body clean of all of them.

Where they persist: in `research/construct_validity_data.csv` (the 2026-08-04 extract, which is the historical dataset and correctly holds historical values), in the superseded manuscript versions v2 through v9, and in the change logs that record the corrections. All are archival.

CURRENT VALUE: 0.623, on 68 labels from 14 raters, paired with intervals computed on that same set.

RISK: GREEN. The one live hazard — a 68-label point estimate printed beside 63-label intervals — was the defect the Final pass repaired. `research/construct_validity_data.csv` is not labelled as superseded in the file itself, which is an archival tidiness matter and not a manuscript risk.

PART 7. JRS CLAIM AUDIT

Claim that must be absent	Present
JRS efficacy	no

Claim that must be absent	Present
JRS superiority	no
JRS validated / validation	no
JRS psychometrically validated	no (heading 8.6 states the negative)
JRS improved reviewer accuracy	no
JRS improved documentation outcomes	no

Current wording, §7: *"For JRS, the result provides preliminary evidence that the record-level distinction embodied in its review logic is operationally detectable; it is not evidence that JRS itself improves documentation outcomes."*

Supporting boundaries retained: §5 states the comparison against unaided judgment is a separate study with its own participants and registration; heading 8.10 states "No criterion validity, and no efficacy."

JRS CLAIM STATUS: PASS.

PART 8. DRR CLAIM AUDIT

Construct	Distinguished	Where
Detectability	claimed, bounded to the constructed corpus	Abstract, §6.1, §7
Reliability	reported and failed against the pre-registered lower bound	§6.5
Criterion validity	explicitly not established	Abstract; §2 table "Not attempted"; 8.10
Construct validity	explicitly construct-dependent	§4.4 "independent of the <i>results</i> and not fully independent of the <i>construct</i> "; §2.4
Psychometric validation	explicitly not established	8.6
Workflow independence	explicitly a design intention, not a result	8.5
Efficacy	explicitly not established	8.10
Generalisability to human-authored records	explicitly not established	8.5
Cross-cultural validity	explicitly not established	8.4
Superiority over unaided judgment	out of scope, separate study	§5

The manuscript's own framing at §9: *"What the paper does not establish is longer than what it does."* The four-rung validity table at §2 marks criterion validity **"Not attempted. Study 4."**

DRR CLAIM STATUS: PASS.

PART 9. REFERENCE-CLASSIFICATION CLAIM AUDIT

Description that must be absent	Present
human raters	no — v3 states "not human raters"
expert raters	no
credentialed professionals	no
expert validators	no
human validation	no — v3 states it "does not constitute independent human validation"
criterion validation	no

Distinction	Stated correctly
Aware of the verification task	YES — "The verification packet identified the task as verification of an answer key"
Access to intended labels	NO — "without access to the intended labels"
Access to author-side classification	NO — "or to the author-side classification"

"Blind" terminology. blind raters and blinded raters appear zero times in v3. The contradicted claim *"were not told that a reference classification existed to be recovered"* is removed. The surviving uses of "blind" refer to the reviewer panel, which was genuinely blind to the reference classification, and to the author-side classification being withheld.

REFERENCE-CLASSIFICATION CLAIM STATUS: PASS.

PART 10. REPRODUCIBILITY DISCLOSURE

Element	Status	Notes
Operational classification rule	FULLY REPRODUCIBLE	research/AnswerKey_Verification_Packet.md, complete text
Intended labels	FULLY REPRODUCIBLE	research/Intended_Key_authorside.md, all 24 with rationale
24-record corpus	FULLY REPRODUCIBLE	released under the data-availability terms in §10
Verified key	FULLY REPRODUCIBLE	research/Verified_Key.md, all 24 labels
Detection-panel results	PARTIALLY REPRODUCIBLE	aggregates in closed_aggregates_2026-08-15.json; producer recoverable only from git history; per-read data RLS-locked
Reliability results	FULLY REPRODUCIBLE	live bench_labels is anon-readable; compute_ac1_ci.py + recompute_current_ac1.py regenerate every value
72 LLM classifications	NOT REPRODUCIBLE	aggregate only; no per-pass outputs
LLM model identity	NOT REPRODUCIBLE	not retained
LLM version	NOT REPRODUCIBLE	not retained
Execution date	NOT REPRODUCIBLE	not retained
Prompt	PARTIALLY REPRODUCIBLE	the rule given to raters is retained; no system prompt or wrapper
Temperature / configuration	NOT REPRODUCIBLE	not retained
Per-pass outputs	NOT REPRODUCIBLE	not retained

Does the missing information materially affect interpretation? Partly. A replicator can apply the same rule to the same corpus with their own raters, human or automated, and check whether the key reproduces — which is the substantive replication. What they cannot do is reproduce *this particular* automated pass. v3 says exactly that.

None of the missing fields can reasonably be recovered. The pass was executed before the disclosure standard was set and the outputs were not written to disk. No recommendation to recover them is made.

PART 11. ETHICS / PARTICIPANT DISCLOSURE

Item	Manuscript statement	Assessment
IRB review	§4.8: <i>"This study was not reviewed by an institutional review board. It was conducted outside any institution holding an IRB, by an independent researcher, with adult professional volunteers"</i>	ADEQUATE. Stated plainly, no exemption claimed or implied
Informed consent	§4.8: participation voluntary, in a personal capacity, withdrawable before publication; attribution opt-in with consent	ADEQUATE
Voluntary participation	§4.8, and Acknowledgments <i>"All 58 worked unpaid, in a personal capacity"</i>	ADEQUATE
Compensation	uncompensated, stated twice	ADEQUATE
Professional identity	contributors named only with consent; may participate anonymously	ADEQUATE
Data retention	responses stored append-only, used only in aggregate	ADEQUATE
Anonymity / confidentiality	records constructed and de-identified; no real case, individual or organisation	ADEQUATE
Withdrawal exercised	§4.8 and Acknowledgments record one panel member who withdrew consent to be named; judgments retained unnamed at her election	ADEQUATE, and unusually candid

No IRB exemption or approval is invented anywhere. The absence is disclosed as an absence.

Assessment: ADEQUATE.

PART 12. NOVELTY CLAIM

CURRENT NOVELTY CLAIM (§2):

"The claim of novelty is narrow and we state it narrowly. It is not that reconstructability has never been valued: it is presupposed throughout administrative law and audit practice. It is that reconstructability of the individual record has not been operationalised as a measurable property with a stated instrument, a stated scale, and reported detection and agreement statistics. That is what this paper supplies..."

SOURCE SUPPORT: none, and none is possible. This is a universal negative over the literature. The repository contains no systematic review, no search protocol and no coverage claim that could support it.

OVERBROAD: YES.

The paragraph does good work: it pre-empts the obvious objection by conceding that reconstructability is long-valued, and narrows the claim to operationalisation. But the narrowed claim is still stated as fact about the whole literature, without qualification. Reviewers in AI governance and measurement routinely require "to our knowledge" on exactly this construction, and its absence invites a reviewer to supply a counterexample and treat the omission as overreach.

CLASSIFICATION: SURGICAL EDIT.

A qualifying clause — "to our knowledge", "we are not aware of", or "we have not found" — resolves it without touching the substance. Not a blocker in the sense of invalidating anything, but it is the single item most likely to draw an avoidable reviewer objection, and it costs three words.

PART 13. CONSTRUCT-DEPENDENCE

Limitation	Present	Location
Author-generated corpus	YES	§4.2 provenance, §2.4, 8.1
Author-defined operationalisation	YES	§4.4 "the reference classification encodes that operationalisation"
Bimodal spectrum / ends of the severity range	YES	§4.3, §7, and the §4.4 concession that unanimity indicates easy cases
Lack of naturalistic records	YES	8.5, "establishes detectability on AI-generated records"
Lack of independent human validation	YES, new in v3	§4.4, "does not constitute independent human validation ... No human replication ... has been performed"
Lack of criterion validity	YES	Abstract, §2 table, 8.10
Investigator dependence	YES	§9, "None of them removes investigator dependence", and the recommendation of an independent adjudicator
Competing interests	YES	§9, first author created the construct and would benefit from adoption

The §2 four-rung table makes the position visible at a glance rather than burying it in a limitations section, and §9 states that what the paper does not establish is longer than what it does.

CONSTRUCT-VALIDITY DISCLOSURE: ADEQUATE.

PART 14. EDITORIAL OVER-DEFENSIVENESS

Qualitative assessment: A, with a trace of C.

25 distinct limitation markers appear across the manuscript. For a paper whose first author created the construct under study, that density is defensible and is itself an integrity signal. A reviewer is far more likely to complain about an overclaiming paper than this one.

Five locations where the same concession is made more than once and the repetition is noticeable:

1. **§4.4 and 8.1** both make the construct-dependence point about the corpus instantiating the authors' operationalisation.
2. **§4.4 "a corpus on which automated raters never disagree is a corpus of easy cases"** and **§4.3 spectrum restriction** state the same easy-corpus concern from two directions.
3. **§6.3 and 8.9** both make the reviewer-heterogeneity point about not licensing individual-level reliance.
4. **§7 and 9** both state that what the paper does not establish exceeds what it does.
5. **Abstract, §2 table, and 8.10** each separately disclaim criterion validity.

Item 5 is arguably correct practice: an abstract-level disclaimer, a structural table and a limitations entry serve different readers.

CLASSIFICATION: OPTIONAL. No edit recommended. Compressing these would trade a real integrity signal for a marginal gain in concision, and every one of the repetitions is accurate.

PART 15. UNRELATED TECHNICAL DEBT

Issue	Affects this manuscript	Classification
<code>compute_ac1_ci.py</code> crash at <code>per_condition_ac1</code>	YES, reproducibility only; no reported value passes through it	REPRODUCIBILITY LIMITATION
<code>compute_ac1_ci.py dedup_last()</code> no-op on the committed CSV	YES, reproducibility only; the 113→104 step is not demonstrable from a committed file	REPRODUCIBILITY LIMITATION
<code>FULL_DATA_ANALYSIS_2026-08-15.txt</code> producer uncommitted	YES, partially; mitigated by <code>recompute_current_ac1.py</code> , which regenerates the reliability block	REPRODUCIBILITY LIMITATION
<code>api/pstat-4c8e1b6a2d90.js</code> deleted, recoverable from git only	YES, partially; the detection-panel producer	REPRODUCIBILITY LIMITATION
<code>construct_validity_data.csv</code> not labelled as superseded	NO	UNRELATED TECHNICAL DEBT
E-11 unresolved identity	NO — no reported figure depends on it	UNRELATED TECHNICAL DEBT
Website deployment	NO	UNRELATED TECHNICAL DEBT
Cloudflare Workers check failing	NO — dashboard-side integration, no repository artifact	UNRELATED TECHNICAL DEBT
Outreach messages, due 31 August 2026	NO	UNRELATED TECHNICAL DEBT
Arm B paper unwritten	NO — separate paper, correctly out of scope here	OPTIONAL FUTURE STRENGTHENING
Commercial pages paused	NO	UNRELATED TECHNICAL DEBT
<code>api/register.js</code> + <code>reference/index.html</code> undeployed	NO	UNRELATED TECHNICAL DEBT
<code>terms.html</code> placeholder spans	NO	UNRELATED TECHNICAL DEBT

PART 16. FINAL SUBMISSION GATE

Issue	Evidence	Status	Submission Impact	Action
Novelty claim is an unhedged universal negative	§2, "has not been operationalised as a measurable property..."; no supporting source exists or could	AMBER	EDITORIAL RISK — most likely avoidable reviewer objection	Add "to our knowledge" or equivalent. One surgical edit
"before any reviewer was recruited" not verifiable	§4.2; <code>pilot_progress</code> exposes no registration timestamp; key committed 2026-07-06, reading by 2026-06-28	AMBER	EDITORIAL RISK	Owner confirms from personal knowledge, or narrow to "before verification began"
72 reference classifications not retained per pass	<code>Verified_Key.md</code> records the aggregate only	AMBER	REPRODUCIBILITY LIMITATION, disclosed in v3	None. Already stated

Issue	Evidence	Status	Submission Impact	Action
Reference LLM identity, version, date, configuration not retained	no artifact	AMBER	REPRODUCIBILITY LIMITATION, disclosed in v3	None. Not recoverable
Section 4.4 vs Appendix A vendor-specificity asymmetry	: 220 names three vendors; §4.4 can name none	AMBER	EDITORIAL RISK — a reviewer may ask	Optional one-clause note; not required
compute_ac1_ci.py crashes as __main__	KeyError: 'basis_identification', line 169	AMBER	REPRODUCIBILITY LIMITATION; no reported value affected	Fix post-submission
dedup_last() no-op on the committed CSV	63 labels both deduplicated and not	AMBER	REPRODUCIBILITY LIMITATION	Fix post-submission
Detection producer recoverable only from git history	closed_aggregates...json names api/pstat-4c8e1b6a2d90.js, deleted at 120c11e	AMBER	REPRODUCIBILITY LIMITATION	Restore or archive post-submission
Reference raters described as automated, not human	v3 §4.4	GREEN	resolved in v3	None
Contradicted blinding claim removed	v3 §4.4	GREEN	resolved in v3	None
Pre-registration deviation, 2 passes vs 3, disclosed	v3 §4.4	GREEN	resolved in v3	None
Judgment denominator 72 stated	v3 §4.4	GREEN	resolved in v3	None
"Nothing withheld" claim corrected	v3 §4.4 Disclosure	GREEN	resolved in v3	None
Human-validation limitation added	v3 §4.4	GREEN	resolved in v3	None
Primary detection statistics trace to source	Part 4 above, 7 of 7 verified	GREEN	none	None
Reliability statistics trace to source	Part 5 above, 9 of 9 verified	GREEN	none	None
No superseded value in the active manuscript	29 checked, body clean	GREEN	none	None
Participant architecture, 58 / 61 / three-person overlap	Part 3 above	GREEN	none	None
Arm A / Arm B expertise parity and disjointness	Part 3 above	GREEN	none	None
JRS claim boundary	Part 7 above	GREEN	none	None
DRR claim boundary	Part 8 above	GREEN	none	None
Ethics and consent disclosure	Part 11 above	GREEN	none	None
Construct-dependence disclosure	Part 13 above	GREEN	none	None
Pre-registration timing requirement met	analysis 2026-08-15, key 2026-07-06	GREEN	none	None
Editorial repetition at five locations	Part 14 above	GREEN	none	None. Integrity signal outweighs concision
Human replication of the answer key	JRS_Validation_Report.md:297 open item 4; blind packet ready	OPTIONAL	would materially strengthen §4.4	Future work
Arm B comparison paper	separate registration	OPTIONAL	out of scope	Future work
construct_validity_data.csv unlabelled as superseded	archival tidiness	OPTIONAL	none	Post-submission
E-11 unresolved identity	no reported figure depends on it	OPTIONAL	none	Post-submission
Website, Cloudflare, outreach, commercial pages, API deployment, terms.html	Part 15	GREEN	none	Unrelated technical debt

SUBMISSION BLOCKERS: 0

AMBER ISSUES: 8

GREEN ITEMS: 21

OPTIONAL FUTURE STRENGTHENING: 4

FINAL RECOMMENDATION

ONE SURGICAL CORRECTION THEN SUBMIT.

The correction is the novelty hedge at §2. It is the only item that is both a manuscript defect and cheap to fix: a universal negative about the literature, stated without qualification, in a paper that is otherwise scrupulous about the limits of what it claims. Three words resolve it.

Everything else that remains is either disclosed in the manuscript already, or is post-submission housekeeping, or is future research. The reproducibility gaps around the automated reference pass are real, are not recoverable, and are now stated in the paper rather than papered over — which is the correct disposition for a limitation that cannot be cured.

The "before any reviewer was recruited" claim needs your confirmation rather than my edit. If you can date the author-side classification against the first recruitment from your own records, leave it. If you cannot, narrowing it to "before verification began" costs nothing and is fully supported by `Intended_Key_authorside.md`.

FINAL STATUS: AMBER

RECOMMENDATION: ONE SURGICAL CORRECTION THEN SUBMIT

MANUSCRIPT MODIFIED: NO

DATA MODIFIED: NO

SCRIPTS MODIFIED: NO

PREREGISTRATION MODIFIED: NO

COMMITTS CREATED: NO