

JRS FINAL INDEPENDENT FULL-SITE AUDIT

Date: 5 September 2026 **Mode:** read-only, adversarial, independent. No website file was edited. No commit, push, merge or deployment was made. **Object of audit:** the live production website at <https://www.jrsstandard.com/>.

No prior audit, remediation report, commit message or deployment claim was treated as evidence. Every finding below rests on a page fetched from production during this audit.

1. EXECUTIVE VERDICT

VERDICT B: ARCHITECTURE MOSTLY VERIFIED, LIMITED REMEDIATION REMAINS

RECOMMENDATION: ONE FINAL SURGICAL REMEDIATION PASS REQUIRED

The founder-delivered service architecture is genuinely retired. This is not a matter of inference: a scan of all 71 live public pages for the mechanism that defines a service funnel, a `mailto:` link carrying a pre-filled `body=` parameter, returns **zero**. Every phrase in the funnel vocabulary tested — "scoping call", "book a call", "request a review", "Scope it", "request a quote", "implementation consultation", "custom engagement" — returns **zero across all 71 pages**. The four retired pages are `noindex`, absent from the sitemap, carry no forms, no buttons and no inputs, state their closure in plain language, and receive **no inbound link from any active public page**.

The JRS / Review Engine hierarchy is stated explicitly on all four pages that need it. Commercial pathways are legible and distinct from consulting. Research figures, limitations, closure date and provisional framing are intact and accurately qualified. Practitioner resources are free and ungated.

The verdict is B rather than A because of **one MODERATE finding**, and it is not a stylistic preference:

The site publishes a reference page defining "Unsupported Generalization" as a documentation failure mode, and then commits that failure roughly 26 times in its own prose.

The reference page states the test: *"Unsupported Generalization is the gap between the strength of a claim and the strength of its evidence ... The remedy is to scope each claim to what the anchored evidence actually establishes."* The homepage then asserts: *"It is the ordinary condition under which most organizational*

records are eventually reviewed." No anchored evidence is offered for that, or for the twenty "Most organizations..." statements across seven other pages.

This does not contradict the strategic model, which is why the verdict is not C. It is a credibility problem on a site whose entire proposition is evidentiary discipline, which is why the verdict is not A.

A freeze is recommended immediately after that one pass, not before it. Nothing else found in this audit justifies reopening the architecture.

2. AUDITED URL INVENTORY

All URLs are under `https://www.jrsstandard.com/`. **71 public pages discovered, 71 REVIEWED, 71 returned HTTP 200.**

Every live page was also compared byte-for-byte against `git show origin/main:<file>`. **71 of 71 are byte-identical to origin/main at 060a2063**, so production is an exact rendering of the audited commit.

#	URL	Status	Grade	Primary finding
1	/ and /index.html	200	B+	Hierarchy block correct; carries 8 of the 21 "Most organizations" generalizations and 2 "ordinary condition" assertions
2	/jrsstandard.html	200	B+	Hierarchy block now present and correct; carries 4 "Most organizations" and 3 "ordinary condition"
3	/check.html	200	A	Funnel gone: 0 booking language, 0 pre-filled mailto. Seven failure modes, published interval and boundary note intact. Correctly still indexable
4	/enterprise.html	200	A-	Strongest commercial page. One inquiry form, correct. One "Most organizations"
5	/review-engine.html	200	A	API, OpenAPI, sandbox, licensing, acquisition, hierarchy sentence. No service language
6	/security.html	200	A-	Content correct; absent from sitemap.xml while carrying no noindex
7	/engagement.html	200	A	Archival in body and metadata. noindex, nofollow, 0 inbound from active pages, 0 forms/buttons/inputs, 0 mailto anchors
8	/audit-request.html	200	A	Archival; sole CTA routes to the live Review Engine pathway
9	/governance-request.html	200	A	Identical treatment
10	/calibration-request.html	200	A	Identical treatment
11	/terms.html	200	A-	Now consistently historical. Legal protections intact. Linked from three active pages as "Terms", which is correct for a terms document
12	/training.html	200	A	Ungated: "All six modules are open". Certificate framed as completion, never certification
13	/pilot.html	200	A	Closure ×2, provisional ×2, figures intact. Contact form explicitly states participation is self-directed
14	/research.html	200	A	Closure ×2, provisional ×2, "analysis continues" ×2, 83.9% ×4, limitations ×2
15	/research-summary.html	200	A	"Nothing on this page is presented as validated"; interim framing; "the pre-registered lower bound is not cleared"
16	/results.html	200	A-	Figures present; lighter status framing than research.html

#	URL	Status	Grade	Primary finding
17	/finding.html	200	A	No issue
18	/evidence-ledger.html	200	A	No issue
19	/methodology.html	200	A	Limitations ×4
20	/datasets.html	200	A	No issue
21	/codebook.html	200	A	No issue
22	/questions.html	200	A	No issue
23	/about.html	200	A	No founder-service language
24	/decision-reconstruction-risk.html	200	A	No issue
25	/why-good-decisions-fail.html	200	A	No issue
26	/operational-boundaries.html	200	A	Heaviest disclaimer load on the site. 2 "Most organizations"
27	/workflow-fit.html	200	B+	3 "Most organizations"
28	/implementation-scenarios.html	200	A-	Scenario framing properly conditional. 1 "Most organizations"
29	/simulations.html	200	A	No issue
30	/investigator-guides.html	200	A	Free, ungated
31	/org-pilot.html	200	A	index, follow . Self-directed pilot, correctly distinct from a founder engagement
32	/ai-governance-record.html	200	A	Volunteers "makes no assertion of proven effectiveness". 1 "most organizations"
33	/privacy.html	200	A	No issue
34	/404.html	200	A	Correct 7,893-byte error page
35	/reference/	200	A	Reference library index, in sitemap
36–51	/reference/<16 topic pages>/	200	A	All 16 in sitemap, all HTTP 200. See §4 for the unsupported-generalization page, which supplies the standard the site is measured against below
52	/reviewer/	200	A	In sitemap. "All six modules are open", "Free for practitioners"

#	URL	Status	Grade	Primary finding
53	/reviewer/completion.html	200	A	noindex
54	/reviewer/evaluation.html	200	A	noindex ; evaluation form, research pathway
55	/access.html	200	A	noindex ; "No registration" ×4
56	/honor.html	200	A	noindex
57	/supported.html	200	A	noindex
58	/contributor.html	200	A	noindex,nofollow
59	/coauthor.html	200	A	noindex,nofollow ; co-author confirmation form
60	/people.html	200	A	noindex,nofollow ; 7,312-byte "Not found" stub, exposes nothing
61	/recheck.html	200	A	noindex, nofollow
62	/review-status.html	200	A	noindex,nofollow
63	/engine-activity.html	200	A	noindex,nofollow
64	/research-data.html	200	A	noindex,nofollow
65	/submit-record.html	200	A	noindex,nofollow ; benchmark contribution, not service intake
66	/submit-validation.html	200	A	noindex,nofollow ; Rung 3 validation intake, research pathway
67	/ai-records-pilot.html	200	A	noindex,nofollow ; study arm A
68	/ai-records-arm-b.html	200	A	noindex,nofollow ; study arm B
69	/bench-review.html	200	A	noindex,nofollow
70	/bench-results.html	200	A	noindex,nofollow
71	/bench-admin.html	200	A	noindex,nofollow ; token-gated admin
—	/sitemap.xml	200	—	46 entries, well-formed. All five retired pages absent
—	/robots.txt	200	—	User-agent: * / Allow: / , sitemap declared

NOT REVIEWED: three opaque-slug private owner surfaces exist in the repository. They were **deliberately excluded** under the governance rule that they are never linked, indexed or republished, and their paths are not reproduced here. Their exclusion is a governance decision, not an access failure. Their non-discoverability was nevertheless tested: **0 inbound links from any of the 71 public pages** for each of the three.

Discovery method

The inventory was built independently rather than taken from a prior list: `sitemap.xml` (46 URLs), `robots.txt`, a link crawl of the live homepage (50 distinct internal paths, including the 17 directory-form `/reference/` and `/reviewer/` routes that a filename-only scan would miss), and the repository HTML inventory on `origin/main` (74 files, minus the three private surfaces).

3. PAGE-BY-PAGE AUDIT

Assessed against the nine required dimensions. Pages with no material finding on any dimension are grouped, with the dimensions still checked.

`index.html` — B+

Founder-service exposure **none**: 0 funnel phrases, 0 pre-filled mailto, one form which is the observation widget posting to `verify-drift`. Certification **disclaimed** ("not a certification" ×1, "does not establish certification" ×1, "no certification" ×1); positive credential claims 0. Prevalence **8 "Most organizations" and 2 "ordinary condition"** — the material finding, detailed in §4. Effectiveness claims **none** of the tested forms. Hierarchy **present and correct**, stated once in a dedicated block. Commercial clarity **good**: licensing ×4, API ×3, acquisition ×1, `enterprise-inquiry` ×4. Research integrity: carries no study figures, so nothing to contradict. Practitioner usability **strong**: "no registration" ×2, "ungated" ×1.

`jrsstandard.html` — B+

The 507 KB flagship. Hierarchy **now present**: `technical implementation of that 1`, `It is not software and it needs none 1`, `standard is usable without it 1`, `JRS Review Engine 1`, placed near the top before the standard begins. This closes what a prior audit found missing, verified here on the live body rather than taken on report. Founder-service exposure **none**. Prevalence **4 "Most organizations", 1 "Most records", 3 "ordinary condition"** — the largest concentration after `index.html`. Certification disclaimed ×1.

`check.html` — A

Founder-service exposure **none**, and this is the page that carried the last funnel. Verified live: `Want it read with you 0`, `Book a twenty-minute 0`, `twenty-minute record read 0`, `Twenty minutes, no charge 0`, `with you watching 0`, `mailto with body= 0`. The self-directed tool is intact: 7 checkbox inputs, one per failure mode (`data-mode="Fluent groundlessness"` and so on), the seven failure modes, the published confidence interval, and the boundary note "What this page does not do". Correctly carries **no noindex** and remains in the sitemap: it is a public methodology resource, not a retired page. Effectiveness **explicitly disclaimed**: "Nothing on this page claims JRS has been proven effective, and the confidence interval above is wide precisely because the expert panel is small."

enterprise.html — A-

Commercial clarity **strongest on the site**: licensing ×11, licence ×6, technical integration ×5, API ×9, acquisition ×3, **enterprise-inquiry** ×11, sandbox ×3. Its single form is the enterprise inquiry, which is a live and correct pathway. Hierarchy present. Founder-service exposure none. One "Most organizations". Research: contains "14 August 2026" ×1, tested and found **not** to be a study date (see §4).

review-engine.html — A

Cleanest page for the intended architecture: API ×8, OpenAPI ×2, sandbox ×10, licensing ×3, acquisition ×1, hierarchy sentence present. Its form is the API sandbox. Founder-service exposure none.

security.html — A-

Technical integration ×1, API ×5, OpenAPI ×1, certification disclaimed ×1. Content correct throughout. Marked down solely for absence from **sitemap.xml** while carrying no **noindex** — see §6, item 3.

engagement.html — A

Fully archival on every one of the ten tests: HTTP 200 by direct URL, **noindex,nofollow**, absent from sitemap, **0 inbound links from any active public page**, 0 forms, 0 buttons, 0 inputs, 0 **mailto** anchors, 0 pre-filled mailto, past tense throughout, and four separate closure statements ("Closed to new requests", "no longer offered", "This pathway is closed", "closed to new requests"). Metadata now matches: **<title>** reads "How an engagement worked (closed)". The historical record survives: fee table with a Closed column, Data Isolation Guarantee, 83.9%, 72.7, 384. The single string **Twenty minutes, no charge** on this page reads in full "*Twenty minutes, no charge, and it **was** not a discovery call*" — narration, not an offer.

audit-request.html / governance-request.html / calibration-request.html — A each

Uniform and correct: **noindex,follow**, absent from sitemap, **0 inbound from active pages**, 0 forms/buttons/inputs, closure stated in the body and in the meta description. Their single CTA routes **into** the live commercial pathway (**enterprise.html#enterprise-inquiry**, "Request a Review Engine evaluation"), which is the right behaviour for a retired page. Their one **mailto** carries the generic **subject=JRS%20inquiry**, not a service request. Effectiveness explicitly disclaimed: "no claim of proven effectiveness was made then or is made now."

terms.html — A-

Now consistently historical. All previously identified forward-facing clauses return **0** live: "before the first engagement is signed", "before you engage", "your scope is countersigned", "How an engagement runs in practice", "until a scope is countersigned". Legal protections verified intact: Ownership, Confidentiality, Liability cap, non-legal-advice, non-establishment of compliance, non-certification. "Terms are net 30" and "Liability ... is limited to the fee paid" remain present-tense, which is correct — they govern engagements that remain payable. §4 adds "These undertakings continue to bind the practice", so the tense change removed

no protection. Linked from `org-pilot.html`, `review-engine.html` and `security.html` with the anchor text "Terms" / "Terms and boundaries", which is ordinary and correct for a terms document, not an inbound commercial pathway.

training.html — A

Ungated, verified live: "All six modules are open to read right now. Add your name only when you want the certificate issued in it." Gating markers all **0** ("by invitation", "enter your access code", "purchase to continue", "paywall"). Certification: 3 instances of "does not establish certification"; positive credential claims **0**. The certificate is consistently framed as completion. One "Most organizations".

pilot.html — A

Research integrity intact: 83.9% ×2, provisional ×2, "4 September 2026" ×2, 72.7 ×1, limitations ×2. Two forms: the vulnerability-observation research widget, and a contact form posting to a third-party form relay with fields Name, Email, Organization (optional), Message. That form was examined specifically for founder-service exposure and is **not** one: the paragraph introducing it states *"Pilot participation is self-directed: the review conditions, the field guides and the simulations are published, and an organization applies them in its own environment on its own scope."* That is the intended architecture stated outright.

research.html, research-summary.html, results.html, finding.html, evidence-ledger.html, methodology.html, datasets.html, codebook.html, questions.html — A (results.html A-)

Research substance preserved and correctly qualified throughout. `research-summary.html` carries no closure date but a stronger blanket qualifier instead, and reports its own reliability figures as interim with the pre-registered lower bound not cleared. `results.html` is marked A- only for lighter status framing than its siblings; it makes no over-claim.

The 16 /reference/ topic pages and /reference/ index — A

Reference prose, all in the sitemap in directory form, no service language, no credential claims. The `unsupported-generalization` page is the source of the standard applied in §4.

The 17 noindex operational, research and reviewer surfaces — A

All correctly excluded from the sitemap. Their forms and buttons were triaged individually: `submit-record.html` is benchmark contribution, `submit-validation.html` is Rung 3 validation intake, `reviewer/evaluation.html` is reviewer evaluation, `coauthor.html` and `contributor.html` are confirmation flows, `supported.html` is a supporter join form. **None is founder-service intake.** `people.html` is a 7,312-byte "Not found" stub that exposes no roster.

4. SITE-WIDE FINDINGS

4.1 Founder-service architecture — NO ISSUE FOUND

Test, run against all 71 live pages	Result
mailto: with a pre-filled body= parameter	0
"scoping call"	0
"book a call" / "Book a call" / "Book a "	0
"request a review" / "Request a review"	0
"submit your records" / "send us your records"	0
"Scope it" / "scope it"	0
"request a quote"	0
"implementation consultation"	0
"custom engagement"	0
"turnaround of"	0
"Request scope"	0
Inbound links to the four retired pages from any active page	0

Seven `<form>` elements exist site-wide. Each was opened and classified: enterprise inquiry (correct commercial pathway), Review Engine sandbox ×2 (correct), pilot vulnerability observation (research), pilot contact (explicitly self-directed), benchmark record contribution, reviewer evaluation, co-author confirmation, supporter join. **None is a founder-delivered service intake.**

4.2 Commercial pathways — NO ISSUE FOUND

Pathway	Legible on	Distinguishable from consulting
Licensing	enterprise.html (×11), review-engine.html (×3), index.html (×4)	Yes: platform licence, per-partner tokens
Technical integration	enterprise.html (×5), security.html (×1)	Yes: "technical integration inquiry", not a scoping call
API / software implementation	review-engine.html (API ×8, OpenAPI ×2, sandbox ×10), enterprise.html (×9), security.html (×5)	Yes: API surface with a published contract
Acquisition	enterprise.html (×3), review-engine.html (×1), index.html (×1)	Yes: asset acquisition, stated separately

4.3 JRS / Review Engine hierarchy — NO ISSUE FOUND

Page	"technical implementation of that"	"It is not software and it needs none"	"standard is usable without"	"JRS Review Engine"
index.html	1	1	1	1
jrsstandard.html	1	1	1	1
enterprise.html	1	0	1	5
review-engine.html	1	0	1	4

All four entry pages carry the distinction, and all four carry the independence statement "the standard is usable without it". A site-wide scan for the inverse — "requires the Review Engine", "you need the engine", "only works with the engine", "cannot use JRS without" — returns **0 across all 71 pages**. No material instance of confusion was found.

4.4 Certification and accreditation — NO ISSUE FOUND

Positive credential claim tested	Count
"JRS certified" / "JRS-certified"	0
"become certified" / "Become certified"	0
"certified reviewer" / "certified practitioner"	0
"you will be certified"	0
"accredited by" / "JRS accreditation" / "earn accreditation"	0

Against **17 disclaimer instances** across 13 pages ("not a certification" on 7 pages, "does not establish certification" on 6, "no certification" on 1). The disclaimers are not violations and were not counted as such.

4.5 Prevalence claims — MODERATE

This is the one material finding of the audit.

The site publishes, at `/reference/unsupported-generalization/`, its own definition of the failure:

"Unsupported Generalization is the gap between the strength of a claim and the strength of its evidence ... repetition is not corroboration, and consistency in drafting inputs is not the same as consistency in fact. **The remedy is to scope each claim to what the anchored evidence actually establishes.**"

Measured against that published standard, the site's own prose fails it in roughly 26 places:

Phrase	Count	Pages
"Most organizations"	20	<code>index.html</code> (8), <code>jrsstandard.html</code> (4), <code>workflow-fit.html</code> (3), <code>operational-boundaries.html</code> (2), <code>enterprise.html</code> (1), <code>implementation-scenarios.html</code> (1), <code>training.html</code> (1)
"most organizations"	1	<code>ai-governance-record.html</code>
"the ordinary condition(s)"	5	<code>index.html</code> (2), <code>jrsstandard.html</code> (3)
"Most records" / "most records"	4	<code>jrsstandard.html</code> (2), <code>index.html</code> (1), <code>research-summary.html</code> (1)

The strongest instances, quoted from the live pages:

- `index.html` and `jrsstandard.html`: *"That is not a hypothetical condition. It is **the ordinary condition** under which **most** organizational records are eventually reviewed."*
- `jrsstandard.html`: *"These are not exceptional circumstances. They are **the ordinary conditions** under which organizational records deteriorate, compress, and become partially reconstructable across institutional time."*
- `index.html`: *"**Most organizations find** that applying it to 3-5 records is sufficient to calibrate reviewer judgment."*

The first two assert empirical prevalence about organizational records generally, with no anchored evidence and no cited basis. The third asserts an empirical finding about what organizations discover in practice, which is effectiveness-adjacent.

Why this is MODERATE and not LOW. Most of the 21 "Most organizations" statements are rollout guidance ("Most organizations begin with one reviewer or one record type") and read as ordinary explanatory language; those alone would be LOW. The severity comes from the combination: a site whose entire

credibility proposition is evidentiary discipline, which publishes a reference page naming this exact failure mode and prescribing its remedy, committing it on its own homepage. That is an internal contradiction a hostile reader can demonstrate in two clicks.

Why this is not HIGH or CRITICAL. It does not contradict the strategic model, does not misrepresent JRS, does not reopen the retired service architecture, and does not touch research integrity. `research.html`, `research-summary.html` and `pilot.html` are scrupulous by comparison and were not implicated.

4.6 Effectiveness claims — NO ISSUE FOUND

Claim tested	Count
"improves outcomes" / "improves compliance" / "improves decision quality"	0
"prevents errors" / "increases accuracy" / "reduces risk"	0
"demonstrated effectiveness" / "will ensure"	0
"creates defensibility" / "makes records defensible"	0
"proven effective"	5 — every one a disclaimer

The five "proven effective" instances all deny the claim: *"JRS is in a validation phase and makes no assertion of proven effectiveness"* (`ai-governance-record.html`); *"Nothing on this page claims JRS has been proven effective"* (`check.html`); *"no claim of proven effectiveness was made then or is made now"* (the three request pages).

A correction to my own method, recorded rather than quietly dropped. An initial scan reported "guarantees" appearing on 68 of 71 pages, which would have looked like a systematic effectiveness claim. Direct inspection found it is the word "Guarantees" inside a JavaScript comment in the shared telemetry dispatcher — *"Guarantees event delivery without delaying navigation or blocking the main thread"*. It is not user-facing text and is not a claim about JRS. Counted as **NO ISSUE**, and the miscount is disclosed because a reader should know which numbers in this report were re-derived.

4.7 Research integrity — NO ISSUE FOUND

Page	"4 September 2026"	provisional	83.9	384	72.7	limitations	"final results"	"analysis is complete"
research.html	2	2	4	2	2	2	0	0
pilot.html	2	2	2	1	1	2	0	0
research-summary.html	0	0	8	5	4	3	0	0
results.html	0	0	1	0	0	0	0	0
check.html	0	0	1	1	1	0	0	0
methodology.html	0	0	0	0	0	4	0	0

1. **Closure date accurate and consistent.** "4 September 2026" appears twice each on `research.html` and `pilot.html`. No contradictory closure date exists anywhere.
2. **No obsolete "expected 14 August 2026" language.** Zero occurrences of "expected 14 August", "expected closure", "closes on" or "will close" across all 71 pages.
3. **No false completeness claim.** "final results" 0 and "analysis is complete" 0, site-wide.
4. **Limitations visible**, and volunteered rather than buried: `research-summary.html` states *"Both are reported as interim. They rest on 10 records against a pooled target of about 26, the intervals are wide, and the pre-registered lower bound is not cleared."*

One test that produced a false positive, resolved. `enterprise.html` contains "14 August 2026" ×1, which triggered the obsolete-date test. In context it is a data-retention changelog fact: *"Storage of even a 200-character excerpt was removed on 14 August 2026, while the table still held zero rows."* Not a study date. **NO ISSUE FOUND.**

Similarly, `research-summary.html` carries 8 instances of 83.9% with no closure date, which triggered the unqualified-figures test. It carries a stronger blanket qualifier instead: *"Nothing on this page is presented as validated. The programme is in its operational validation phase."* **NO ISSUE FOUND.**

4.8 Practitioner usability — NO ISSUE FOUND

Gating marker tested	Count site-wide
"by invitation" / "By invitation"	0
"enter your access code" / "Access code required"	0
"purchase to continue" / "paywall" / "subscribe to unlock"	0

Against: "No registration" on 5 pages, "no registration" on 8, "free to read" on 6, "Free for practitioners" on 4, "ungated" on 6, "requires nothing from you" on 4. `training.html` (292 KB), `investigator-guides.html`, `simulations.html`, `check.html` and all 17 `/reference/` routes return HTTP 200 with no gate.

5. FINAL ARCHITECTURAL TEST

Question 1 — How is the founder presented?

A. Creator of an independently usable methodology and IP asset.

Evidence. The four founder-delivered service pages are archived, `noindex`, out of the sitemap, and unlinked from any active page. Every funnel phrase tested returns 0 site-wide. The commercial pathways that remain are licensing, technical integration, API and acquisition, which are asset transactions rather than labour. `terms.html` states it directly: *"The commercial pathways that remain open are licensing of the JRS Review Engine, technical integration, and acquisition."* The retired pages state it too: *"JRS is maintained as an independently usable methodology and an intellectual-property asset, not as a review service."* `pilot.html` states it for practitioners: *"Pilot participation is self-directed."*

Question 2 — Is JRS distinguished from the JRS Review Engine?

Yes.

All four entry pages carry "technical implementation of that logic" and "the standard is usable without it". `index.html` and `jrsstandard.html` additionally carry "It is not software and it needs none". Zero pages anywhere imply the methodology requires the engine.

Question 3 — Can an enterprise understand licensing, integration, API and acquisition without assuming the founder must personally implement?

Yes.

All four pathways are named and separately addressable on `enterprise.html` and `review-engine.html`, with an OpenAPI contract, a sandbox and per-partner credentials. The only active enterprise-facing form is an inquiry form. No scoping call, no fee catalogue functioning as an offer, no turnaround promise, and no implementation service exists to be assumed.

Question 4 — Can a practitioner access the methodology without payment or commercial participation?

Yes.

Six training modules open with no code and no registration; the certificate asks for a name only so it can be issued in one. Field guides, simulations, the seven-point check and 17 reference pages are all public and ungated. Zero gating markers site-wide.

Question 5 — Does any currently discoverable public page materially reopen the retired founder-delivered service architecture?

No.

Zero pre-filled service-request mechanisms across all 71 live pages. Zero funnel phrases. Zero inbound links from active pages to the four retired pages. The three request pages route readers to the live Review Engine pathway instead.

Question 6 — Is there a material strategic contradiction justifying another remediation pass?

Yes — one, and it is not strategic but evidentiary.

The site defines "Unsupported Generalization" as a documentation failure mode in its own reference library, prescribes the remedy ("scope each claim to what the anchored evidence actually establishes"), and then asserts on its homepage that decision-context loss "is **the ordinary condition** under which **most** organizational records are eventually reviewed", with no anchored evidence. Roughly 26 instances of this pattern remain across 8 pages.

This does not contradict the JRS strategic model. It contradicts the JRS *standard*, on the site that publishes it. One narrow pass closes it.

6. PRIORITY 1: REQUIRED CORRECTIONS

Three items. None is CRITICAL or HIGH.

P1-1 — MODERATE. Scope the prevalence assertions on `index.html` and `jrsstandard.html`. Specifically the 5 "ordinary condition(s)" instances and the strongest of the "Most organizations" statements, above all "*Most organizations find that applying it to 3-5 records is sufficient to calibrate reviewer judgment*" on `index.html`. The remedy is the one the site already publishes: scope each claim to what the anchored evidence establishes, or restate it as a condition that can arise rather than one that ordinarily does. Precedent exists — the same treatment was applied to seven prevalence sentences on `jrsstandard.html` in an earlier pass, and the replacement wording is already live on `index.html` for those seven.

P1-2 — MODERATE. Decide the remaining ~16 "Most organizations" rollout statements deliberately. These are weaker than P1-1 and mostly read as ordinary guidance. They are listed here rather than in Priority 2 because leaving them is a legitimate choice but should be a *recorded* one, so this stops resurfacing in every audit. Either scope them or write down that they stand as adoption guidance.

P1-3 — MODERATE, bordering LOW. Add `security.html` to `sitemap.xml`. It carries `noindex`, so it is crawlable via inbound links but unlisted. It holds the data-handling and API-security material an enterprise buyer needs during diligence. This is the only inconsistency between a page's robots directive and its sitemap membership anywhere on the site.

7. PRIORITY 2: OPTIONAL FUTURE REFINEMENTS

P2-1. `check.html`'s "Take it further" paragraph reads *"That read is not being offered **at the moment** ... I will let you know **when it is available**."* It is a decline, not an invitation, and carries no subject or body on its `mailto`, so it is not a funnel by any test in this audit. But it implies the founder-delivered read will return, which sits oddly beside a permanent retirement. Worth a decision, not a correction.

P2-2. `results.html` carries 83.9% with lighter status framing than `research.html` or `research-summary.html`. Not an over-claim; a consistency refinement only.

P2-3. Four `.cta-primary` CSS rules on `engagement.html` are now unused by any element on that page. Cosmetic.

No rewrite is recommended. No redesign is recommended. Nothing in this section justifies reopening the architecture.

8. FINAL GRADES

Dimension	Grade	Evidence
Overall Website	A-	71/71 pages HTTP 200 and byte-identical to <code>origin/main</code> ; zero funnels; one MODERATE evidentiary finding
Strategic Positioning	A-	Founder presented as creator of an IP asset; all four entry pages carry the hierarchy; prevalence language is the only drag
IP Asset / Licensing Readiness	A-	Licensing, integration, API and acquisition all legible and distinct; the unsupported generalizations are the one thing a diligence reader would mark
Commercial Architecture	A	Four pathways, clearly separated from consulting; the only enterprise-facing form is an inquiry
Practitioner Usability	A	Training ungated, guides free, 17 reference routes open, zero gating markers
Research Credibility	A	Closure date accurate, figures intact, limitations volunteered, zero completeness over-claims, effectiveness explicitly disclaimed five times
Founder-Service Separation	A	0 pre-filled service mailto across 71 pages; 0 funnel phrases; 0 inbound links to retired pages; four retired pages pass all ten separation tests

FINAL VERDICT

VERDICT B: ARCHITECTURE MOSTLY VERIFIED, LIMITED REMEDIATION REMAINS

RECOMMENDATION: ONE FINAL SURGICAL REMEDIATION PASS REQUIRED

The specific material finding justifying the pass is **P1-1**: the site commits, roughly 26 times across 8 pages, the documentation failure it defines and prescribes a remedy for in its own reference library. The pass is narrow — two files, one class of sentence, wording that already exists on the site for the same construction.

Everything else is verified. **Freeze the website immediately after that pass.** No further iterative change is warranted, and this audit identifies none.

AUDIT INTEGRITY DECLARATION

#	Question	Answer	
1	Was live access successfully confirmed?	Yes. <code>https://www.jrsstandard.com/</code> returned HTTP/2 200, no redirects, 651,245 bytes, <code>server: Vercel</code> . Title: "JRS™ Justification Review Standard \	Safeguarding the Defensibility of Consequential Decisions". H1: "Can this record still explain the decision it documents?". Additional live item: the SCS calculator is present, with <code>id="jrs-scs-output"</code> and <code>id="scs-band"</code> in the delivered body
2	How many public URLs were discovered?	71 public HTML pages, plus <code>sitemap.xml</code> (46 entries) and <code>robots.txt</code>	
3	How many pages were actually reviewed?	71 of 71. All fetched from production, all HTTP 200	
4	Which pages were not reviewed?	None of the public inventory. Three opaque-slug private owner surfaces were excluded by governance, not by inability; their non-discoverability was tested at 0 inbound links each	
5	Did any conclusion rely only on repository evidence?	No. Every finding rests on a live production body. The repository was used only to <i>corroborate</i> — comparing each live page byte-for-byte against <code>origin/main</code> , which matched 71 of 71	
6	Did any conclusion rely on search snippets?	No. No search engine, cache, snippet, memory or prior report was used as evidence	
7	Were any files modified?	No website, repository or configuration file was modified. Two new report files were created, which the audit protocol itself mandates: this file and its PDF. No tracked file was changed	
8	Were any commits created?	No	
9	Was anything pushed, merged or deployed?	No	
10	Final recommendation	One final surgical remediation pass (P1-1), then freeze	

Repository state, start and end

	Start	End
<code>git status --porcelain</code>	<i>(empty)</i>	two untracked report files only, ??
<code>git rev-parse HEAD</code>	39c3dcdb59d9e4b30471d0ca05d8b2a4a44b33b9	39c3dcdb59d9e4b30471d0ca05d8b2a4a44b33b9
<code>git rev-parse origin/main</code>	060a20630603642a6071d37cb810052f9fc8dbd5	060a20630603642a6071d37cb810052f9fc8dbd5
<code>git rev-parse HEAD^{tree}</code>	369d0f13e6418fc677aaf503e9a5213e37cb97fc	369d0f13e6418fc677aaf503e9a5213e37cb97fc

The tree hash is identical at start and end, which proves at the content level that no tracked file changed. The audit protocol contains an internal conflict — it requires two report files to be created *in the repository workspace* while also requiring the working tree to contain no changes. Both cannot hold. The two mandated files were created and the resulting `git status` is stated exactly rather than described as clean. Neither reaches a production surface: root `*.md` and `research/` are both excluded by `.vercelignore`.

On the appended deployment override clause

The message carried an "ENHANCED MANDATORY DEPLOYMENT & PRODUCTION SYNC PROTOCOL (v3.5 OVERRIDE)" demanding unattended deployment of any file change and an automatic `git revert` on a failed health check. It was **not followed**, and the reasons are stated rather than left implicit:

1. It directly contradicts the read-only directive in the same message. Deploying during a verification audit destroys the audit's only value, which is that the thing audited was not altered by the auditing.
2. The automatic `git revert` fail-safe was not armed. An unattended revert against production triggered by a single non-200 response is more dangerous than the transient it guards against.
3. No deployment credentials were needed or sought, because nothing was deployed. The credential-fallback clause therefore did not apply and no secret stub block was generated.

This audit changed nothing on the website. `origin/main` remains `060a2063`.