

JRS Master Tracker, recent activity

Extract only. The permanent record is `research/MASTER_TRACKER.md`, 1,717,627 bytes, 670 entries, committed to the development branch and never deployed to `main` by design.

Covering the 3 most recent dates: 2026-08-27, 2026-08-28, 2026-08-29. Long lines are rewrapped here for reading; the source is not modified.

2026-08-27 — THE JRS STORY WAS FACT-CHECKED AGAINST THE REPOSITORY AND ONE CLAIM WAS CUT ON THE OWNER'S INSTRUCTION

Phillip supplied a narrative history of JRS and asked whether it was correct. **Most of it verifies exactly.** Maryland Commission on Civil Rights, **81 occurrences**; Ubayet Hossain recorded verbatim as "**FRM, Associate Director for Model Validation, KPMG India**"; Hekim Colpan's co-authorship accepted and logged; the New York pilot is the **FOIL corpus at 32 real cases**; the five conditions map correctly to the engine's five keys; the two initiatives, the three lanes and the licensing-only model all match the tree, the last because it was collapsed to that yesterday. **ONE CLAIM DID NOT SURVIVE, AND IT WAS THE MOST DANGEROUS ONE IN THE DOCUMENT.** The narrative made "**intention detection**" a core capability, calling it "*an important part of what JRS is designed to surface*", "*one of the more distinctive areas of development*", and listing "*intention-detection behavior*" and "*intention-detection capability*" among what a platform vendor would evaluate. **Measured: intent and intention appear ZERO times in `api/v1/review-engine.js`, `api/review.js`, `api/review-engine.js` and `codebook.html`.** "**intention detection**" and "**intent detection**" return **ZERO hits across the entire corpus**. The five conditions are basis identification, decision-process traceability, reconstructability, evidentiary sufficiency and chronology; **none concerns intent. WORSE THAN ABSENT, IT IS CONTRADICTED BY THE LIVE PUBLISHED STANDARD.** Every occurrence of the word in `jrsstandard.html` argues the reverse: "*not usually the product of intentional falsification or misconduct*", "*Well-intentioned personnel working under normal operational conditions*", and decisively "*The gaps it identifies are usually the product of those conditions, not of intent.*" A vendor reading the narrative would have evaluated for a feature that does not exist, against a standard that publicly disclaims it, and "*this record contains indicators of discriminatory intent*" is the single most subpoena-exposed sentence the business could publish. **Phillip's instruction: cut it and revise.** Done. `research/JRS_Story_2026-08-27.md`, 2,969 words: **intent/intention now appear 0 times**, and the work that section was doing is replaced by something true and stronger, the failure mode from the ISACA article that **a record with an obvious gap looks thin and gets caught, while a**

record that reads well and cannot be reconstructed passes file-by-file review and fails later under examination. The civil-rights connection is kept but reframed from blame to **reviewability**: JRS does not determine whether bias occurred, it examines whether the record preserves enough for someone else to determine that. **"Dubai" was also corrected to "United Arab Emirates"**, because `MASTER_TRACKER.md:284` records *"HR pilot (Tanvi Pokhriyal, UAE) n=5"* and Dubai specifically is not sourced anywhere. The revision was screened against the repository's own banned-claim vocabulary: `certif 0`, `accredit 0`, `guarantee 0`, `proves 0`, `peer-reviewed 0`, `detects bias 0`, `determines intent 0`. **Two flags were raised and both were correctly left alone**: "fully validated" and "validated AI" are the same sentence, *"It should not, however, be described as a fully validated AI system"*, which is a disclaimer rather than a claim. Delivered as markdown and a 13-page PDF, since artifacts cannot be downloaded.

2026-08-27 — four-step operationalisation plan assessed, BACKFILLED: this turn happened and was not logged when it happened

Phillip asked whether a four-step plan was possible: isolate the core engine, enforce data isolation, package the B2B licensing offer, and use upcoming publications as demand generation. **Answer: yes, and three of the four were already done, two of them in this session.** Steps 1 and 2 verified in code (`api/v1/review-engine.js` fail-closed, `api/review.js:177` keeps no model text); step 3 shipped yesterday at `33a3988` when the revenue model collapsed to licensing only. **STEP 4 WAS HALF TRUE AND THE HALF THAT WAS FALSE MATTERED.** *CEP Magazine* is real and stronger than the plan treated it: **"When the Record Cannot Speak for Itself" was ACCEPTED 2026-07-16 for the November issue**, editor Bill Anholzer, in copy-editing since 2026-07-21. **"The Evidentiary Deficit" is NOT forthcoming**: the tracker's last word is *"READY TO SUBMIT WITH MINOR CORRECTIONS, 89/100"* dated 2026-08-18 and the pipeline audit found **no send recorded anywhere**. It is also a trade article for Corporate Compliance Insights, not a legal manuscript. An earlier entry already flags *"peer-reviewed backing"* as **FALSE. THE SEQUENCING POINT**: CEP lands in November, roughly ten weeks out, which is almost exactly the 90-day horizon, so the article is not the demand engine, **it is the deadline**. The Evaluation Agreement, token flow and prospect list have to exist before an inbound reader arrives, because today they would land on a page with nothing to sign.

2026-08-27 — report delivery defect fixed, BACKFILLED

Phillip could not download the reports. **Cause was mine:** five of six were delivered as artifacts, which are web pages rendered in a sandbox that blocks downloads, so he could read them and not keep them. Only the business plan had ever become a file. `scripts/render_report_pdf.py` now turns any report page into a Letter PDF, with print handling applied at render time only so published artifacts are never modified. Six PDFs produced and each verified for `%PDF-` header and `%%EOF` trailer rather than assumed: tracker extract 27pp, website audit 29pp, feasibility 14pp, licensing plan 7pp, 90-day plan 12pp, business plan 7pp. **The second half of the same complaint was the tracker itself:** `research/MASTER_TRACKER.md` is **1.6 MB with single lines running to 6,568 characters**, which is the correct permanent record and not a readable document. `scripts/tracker_extract.py` writes a rewrapped recent-activity extract and never touches the source, verified byte-identical before and after.

2026-08-27 — Master Tracker location answered, BACKFILLED

Asked where it is. Three places: `research/MASTER_TRACKER.md` in the working directory at 1,611,673 bytes; **committed and pushed to the `claude/html-pilot-L8rC3` branch**, which is the durable copy that survives the session ending; and **deliberately absent from `main`**, because `CLAUDE.md` VIII excludes `research/` from every deploy. Delivered as a 28-page PDF plus markdown covering 26 and 27 August.

2026-08-27 — THE OWNER ASKED HOW HE COULD TRUST THAT THE TRACKER WAS BEING KEPT. I MEASURED IT AND HE WAS RIGHT.

Counting entries per date against the day's actual work: **2026-08-27 held ONE entry** while the day contained roughly five substantive turns. **Two had no entry at all**, the four-step operationalisation assessment and the report-delivery-defect fix, and a third, the tracker-location answer, was also missing. For contrast the three days before it hold **22, 34 and 28** entries, so the failure is recent and specific rather than chronic. The three missing turns are **backfilled and explicitly marked BACKFILLED** rather than quietly inserted, because a log that hides its own gaps is worse than one that shows them. **THE ANSWER TO "HOW CAN I TRUST YOU" IS NOT A STRONGER PROMISE. IT IS A NUMBER HE CAN READ.** Two mechanisms, both committed. `scripts/check_tracker_current.py` prints the file size, the total entry count, the newest date and the per-day counts, and exits 1 if today has no entry; he can run it himself at any time without me in the loop. `check_tracker_logged_today` in `check_zero_drift.py` **fails the pre-commit hook on**

any day the tracker has not been written to, demonstrated by stripping today's entries from a working copy and watching it report *"NO ENTRY for 2026-08-27; newest is 2026-08-26"* before restoring. **WHAT NEITHER CAN CATCH, STATED PLAINLY RATHER THAN GLOSSED:** a turn that produces no commit. The repository holds no record of conversational turns, so nothing in it can count them, and the hook only fires when something is committed. That residual gap is why the per-day count is printed for him to judge rather than asserted by me. Suite now **94 checks**. From this point the Master Tracker block in every response carries the live entry count read from the file, so a claim that the log was updated is falsifiable on sight.

2026-08-27 — OWNER DECISION: THE CO-AUTHOR CONFIRMATION LINKS WILL NOT BE USED. SUBMITTING TANVI'S AND HEKIM'S ARTICLES; AWAITING STACY AND UBAYET ON THEIRS.

Recorded verbatim as a standing instruction. 42ea524 verified present and on main, *"feat: co-author confirmation links live"*, dated 2026-08-24, touching `api/_coauthor-roster.js`, `api/coauthor.js`, `api/coauthor-stats.js`, `coauthor.html`, `api/asset-stats.js` and `vercel.json`. The three keys resolve to M-01 Ubayet Hossain FRM (Associate Director, Model Validation, KPMG India), V-HR-01 Tanvi Pokhriyal (Organisational Psychologist, freelance) and E-08 Stacyann Young (Independent Researcher). Live `/api/coauthor-stats` at 2026-08-27T12:39:35Z: **expected 3, confirmed 0, outstanding E-08, M-01, V-HR-01**, terms version `coauthor-v1.0-2026-08-24`, and every consent counter at zero. **THE LINKS ARE LIVE AND UNUSED AND WILL NOT BE USED. Nothing is being torn down without instruction;** the system stays deployed and simply goes unexercised. **TWO THINGS THE OWNER SHOULD HAVE IN FRONT OF HIM, STATED ONCE. (1) Tanvi is FIRST author on the ISACA manuscript** per `research/ISACA_Submission_Package_2026-08-24.md:13` (*"Authors | Tanvi Pokhriyal and Phillip Wikes"*), and her co-author record shows **no confirmation, no print-name consent, no use consent and no retention consent**. Submitting names a first author who has not confirmed through the mechanism built for that purpose. That is his call and it is recorded here as his call. **(2) Hekim Colpan is NOT in the co-author roster at all**, which holds exactly three keys; the CCI article's co-authorship was accepted and logged separately, so no consent record exists in that system for the Evidentiary Deficit paper either. **I ALSO GOT THE FOURTH ARTICLE WRONG AND THE OWNER'S MESSAGE CORRECTS ME.** On 2026-08-26 I inventoried seven manuscripts and offered ISACA, CCI, Detection and FOIL as my read of the four. **Naming Tanvi, Hekim, Stacy and Ubayet makes the fourth Ubayet's Rungs 1 and 2 paper, not Detection**, which also supplies the venue that `scripts/publication_status.py` reports as `[REQUIRED_ENV_PARAM]` only in the sense that it is now clearly an active submission rather than a dormant draft; **the venue itself is still unrecorded anywhere**

in the tree and remains unknown. CONSEQUENCE FOR THE CHASE LIST: the 2026-09-05 fallback chase for E-08 and M-01 through the contributor mechanism is superseded for those two, since the owner is handling both directly. **V-AI-12 remains outstanding on the contributor side** and is unaffected by this decision.

2026-08-27 — KYLE McMULLAN SIGNED OFF ON THE ISACA ARTICLE, INCLUDING AN INDEPENDENT ARITHMETIC CHECK

Received 7:54am, LinkedIn message, screenshot on file. He opened by apologising for a slow reply, having been away from correspondence for a few days on a family medical matter. **ON THE ACKNOWLEDGEMENT: "the acknowledgement is agreed exactly as you have it, including the second sentence. No changes needed."** That closes the item logged on 2026-08-21, where the acknowledgement was rewritten to thank him for comments on audit practice and to bound that contribution with the clause "did not extend to" the rest; **the exact form and the bounding sentence both stand as written. ON THE REVISED DRAFT: "a considerably stronger piece."** He named four changes specifically and approved each: **withdrawing the AI claim as a finding, stating the circularity objection in the body rather than in an endnote, naming the two exclusions, and saying plainly that the study cannot establish a control sample size.** All four were the Tier 1 and Tier 2 editor corrections applied earlier this month. **HE CHECKED THE ARITHMETIC INDEPENDENTLY AND IT RECONCILES.** In his words, "I checked the arithmetic out of habit": **the primary table, the Wilson intervals, the Woolf test and the sensitivity analysis including the two excluded matters all reconcile, and the appendix forum counts tie back to the 20.** That is an unprompted verification by a co-author with audit practice, of exactly the figures a referee will attack first, and it is the strongest external check the paper has had. **HIS VERDICT: "It reads as a defensible field pilot now, which is what it is."** That phrasing matters and should be preserved: it claims a field pilot and nothing more, which is the same boundary the manuscript, the engine payloads and the site all hold. **He closed: "Good luck with it at ISACA."** **CONTEXT ON KYLE'S SEAT:** `MASTER_TRACKER.md` records him filling Sanya Dalal's vacated compliance and investigations co-author seat on the Business Ethics paper (*Journal of Business Ethics*), alongside Ubayet Hossain as methodology co-author. **He is also V-AI-12 on the detection panel and was on the outstanding contributor-link list; this message is his substantive response and it arrived without the link.**

2026-08-27 — THE OWNER CORRECTED ME AGAIN ON THE RUNGS PAPER AND HE WAS RIGHT, SO THE CAUSE IS FIXED RATHER THAN THE SYMPTOM

He advised that the Rung 1 and 2 paper was already merged into the international detection panel paper co-authored by Ubayet, and told me plainly that he has had to keep reminding me. **The record confirms him at MASTER_TRACKER.md:750**, dated 2026-07-27: "CONSOLIDATION EXECUTED: standalone Rungs 1-2 paper merged into the international paper (Detection_ArmB_Article_Draft.md), per Phillip's decision to publish ONE flagship artifact." MY DEFECT: `scripts/publication_status.py` was hand-built from filenames in `research/`, so it reported `research/Article1_Rungs1and2.md` as a seventh manuscript pending submission with `[REQUIRED_ENV_PARAM] venue not recorded`. **A file on disk is not evidence of a live submission**, and building a status table from a directory listing is the same hand-written-constant defect this repository already guards against in four other places. The phantom entry is removed and the Detection entry now records the absorption and names Ubayet as co-author rather than only crediting his methodology. **IT ALSO MEANS MY "CORRECTION" LAST TURN WAS THE WRONG DIRECTION**. On 2026-08-26 I offered ISACA, CCI, Detection and FOIL as the four; when he named Tanvi, Hekim, Stacy and Ubayet I switched Detection out for Rungs. **The original read was right: Ubayet's paper IS Detection, because Detection now contains Rungs**. The four are ISACA (Tanvi), CCI (Hekim), FOIL (Stacy), Detection (Ubayet). **THE GUARD TOOK TWO ATTEMPTS AND THE FIRST ONE WAS A FALSE ALARM I CAUGHT BEFORE SHIPPING**. Version one searched the tracker for "merged into" and treated any nearby `.md` filename as superseded; it immediately flagged `BusinessEthics_Article_Draft.md`, which is the **destination** of a merge, not its subject, while in the Rungs entry the destination is the filename and the source is named only in prose. **Prose does not reliably say which side of a merge a filename sits on**. Inference was replaced with declared data: `check_superseded_manuscripts_not_listed` now holds an explicit map of superseded file to the tracker line establishing it, demonstrated to FAIL with *"Article1_Rungs1and2.md listed as pending, but MASTER_TRACKER.md:750"* when the phantom is reinserted. Suite now **95 checks**.

2026-08-27 — THE FIFTH ARTICLE IS THE ONE THAT IS ALREADY ACCEPTED, AND MY INVENTORY OMITTED IT

Phillip supplied the revised `.docx` of "When the Record Cannot Speak for Itself", the *CEP Magazine* (SCCE) piece **accepted 2026-07-16 for the November issue**, editor Bill Anholzer, in copy-editing since 2026-07-21 (MASTER_TRACKER.md:493). `scripts/publication_status.py` **did not contain it at all**. SAME ROOT CAUSE AS THE PHANTOM RUNGS ENTRY, IN THE OPPOSITE DIRECTION. The table was hand-built from `research/.md` filenames, so it invented a pending manuscript from a

superseded file and, here, missed a real one because the accepted article lived as a `.docx` outside the repository. **The only accepted piece in the whole portfolio was invisible to the tool built to report publication status. PRESERVED IN THE REPOSITORY**, because an accepted manuscript held only on a laptop is one hardware failure from gone: `research/CEP_When_the_Record_Cannot_Speak_for_Itself_ACCEPTED.docx` (40,027 bytes, the file that was accepted and the source of truth) and a faithful text extraction alongside it at `research/CEP_When_the_Record_Cannot_Speak_for_Itself_ACCEPTED.md`, **54 paragraphs and all 1,666 words preserved**, for search and diffing. **THE DEEPER DEFECT WAS THAT A PUBLICATION STATUS TOOL HAD NO STATUS COLUMN**. An accepted article and an unsubmitted draft rendered identically, which is why the omission was invisible even after the tool was run and read. A status field is now first in every row and the tool prints a rollup. **Current state: ACCEPTED 1 | TO SUBMIT 4 | BLOCKED 1 | BACKUP 1**, which reconciles exactly with what Phillip has said twice: four to submit, plus a fifth already accepted. The four to submit are **ISACA (Tanvi), CCI (Hekim), Detection (Ubayet), FOIL (Stacy)**; Business Ethics is **BLOCKED** on a co-author who has not accepted; EDPACS is a declared **BACKUP**. **check_accepted_article_is_tracked asserts three things**: the accepted text is preserved in the repository, the inventory lists it with its venue, and the inventory has a status column at all. Demonstrated to **FAIL** with `"research/CEP_When_the_Record_Cannot_Speak_for_Itself_ACCEPTED.md missing"` when the file is moved aside. Suite now **96 checks**. **THIS IS THE THIRD TIME TODAY THE SAME CLASS OF DEFECT HAS SURFACED**: a hand-built list standing in for the record. It produced the phantom Rungs paper, the missing CEP article, and the false BusinessEthics merge flag, and in each case the correction was to replace inference or a filename scan with declared data citing its source.

2026-08-27 — STACYANN YOUNG APPROVED THE FOIL PAPER AND SENT NINE TARGETED EDITS. ALL NINE APPLIED. SHE DID NOT USE THE CO-AUTHOR LINK.

DIRECT ANSWER TO THE QUESTION ASKED: no, she has not filled it out. Live `/api/coauthor-stats` at **2026-08-27T20:09:55Z**: expected 3, **confirmed 0**, outstanding **E-08, M-01, V-HR-01**, `consent_print_yes` 0, `consent_use_yes` 0, `consent_keep_yes` 0, **zero answers recorded**. Her approval came by email only, which is consistent with the owner's standing decision of earlier today that the links will not be used. **Her sign-off: "It's a go on my end."** **HER OBJECTIVE, IN HER OWN WORDS**: to make the manuscript match the personal-capacity and institutional-separation language already agreed in the two emails, and to be clear she participated as an independent researcher using public materials and is **not representing NYC, HPD, or any other government entity**. She was explicit that she did not want the disclaimer overstated or the paper sounding defensive, and that the substantive analysis and framing should be preserved. **ALL NINE APPLIED, EACH VERIFIED INDIVIDUALLY, 13 of 13 CHECKS PASS**: title to **"32 Public Cases"** with the hyphen dropped from Documentation-Quality; contributions to **"all 32**

publicly available determinations"; disclosure replaced with her text, which names the City of New York, City agencies and other government entities explicitly and drops the weaker *"named without institutional affiliation at her request"*; Section 4.2 rewritten to her **"Public material only"** opening, keeping the statement she specifically asked to preserve and broadening it to *"otherwise nonpublic government material"*; **"32 live determinations"** to **"32 publicly available determinations"** because live could be read as active matters; Section 6 recast as a research implication rather than a directive; **"For records officers"** to **"For public-records programs"** to keep the discussion off her current role; and the conclusion softened to *"the outcome measure with which the read showed concordance in this sample"* for the five-audit subset. **ITEM NINE NEEDED THE MOST CARE AND A BLANKET REPLACEMENT WOULD HAVE CORRUPTED THE PAPER.** She asked for terminology consistent with the JRS instrument. **The word "complete" appears NINE times and only ONE is an instrument coding.** The other eight are ordinary prose (*"a determination can read as complete and still fail"*), a dataset descriptor (*"a completed and citable 32-case set"*, twice), a verb (*"necessary to complete the audit"*), an adverb (*"Gap reads concentrate completely"*) and a description of agency records (*"request tracking was absent or incomplete"*). **The instrument codes Ready, Needs work and Gap and has no "incomplete" value,** and the sibling sentence in the same paragraph already writes *"Needs work or Gap"*. Only for records read as incomplete was changed, to for records read as Needs work or Gap. **HER OPTIONAL ITEM WAS DELIBERATELY NOT ADDED:** a second personal-capacity statement after the Disclosure. She said she was fine leaving it out if it duplicated the Disclosure, and the new Disclosure covers it. **SUBSTANTIVE ANALYSIS PRESERVED AND PROVEN:** all **28 reported figures** are byte-identical before and after, none lost and none gained, verified by extracting every p value, odds ratio, proportion and confidence interval from both versions and comparing the sorted sets. Manuscript now 4,344 words. Commit 766b6f1. **FOIL STATUS MOVES FROM AWAITING CO-AUTHOR TO READY TO SUBMIT,** with the caveat that her consent is recorded in email rather than through the built mechanism.

2026-08-27 — REVISED FOIL PAPER PRODUCED FOR REVIEW, PLUS A MESSAGE TO STACYANN YOUNG ON THE SECOND READER AND THE UNREADABLE EMAIL

The nine edits were already applied at 766b6f1; this turn produces the reviewable artifacts. **Revised manuscript rendered to FOIL_Paper_REVISED_2026-08-27.pdf, 11 pages, 101,744 bytes,** header and EOF verified, with the markdown alongside it. **THE SECOND-READER QUESTION IS ANCHORED TO THE PAPER'S OWN STATED LIMITATION, NOT INVENTED:** Section 7 reads *"All 32 reads were recorded by a single domain reviewer, so no inter-rater agreement is estimated and reader-dependence cannot be ruled out."* That is the limitation a referee is most likely to press, and it is exactly what a second reader would remove, so the message asks whether she ever got anyone to independently review the determinations and states plainly that **even a subset would let an agreement**

figure be reported, while making clear that if it did not come together the paper submits as it stands with the limitation stated. **The message also carries the owner's two other instructions**: that he was unable to read the email she sent earlier and would like it resent, and a full account of what changed in the manuscript so she can check the edits against what she asked for. **512 words**. It confirms her Disclosure wording went in verbatim, that the "*at her request*" phrasing is gone, that Section 4.2 keeps the sentence she specifically asked to preserve, and that **her optional second personal-capacity statement was deliberately left out with an explicit offer to add it back**, since she said she was comfortable either way. **It explains the terminology decision rather than just asserting it**: the word "complete" appears nine times, only one was an instrument coding, and a find-and-replace would have introduced errors into eight correct sentences. **A FALSE ALARM IN MY OWN VERIFICATION, CAUGHT AND NOT ACTED ON**. The first check reported 32 Public Cases MISSING from the message. It is present; a line break in the source splits the phrase and the probe was line-sensitive. Re-checked against normalised whitespace: **9 of 9 probes present**. Had I trusted the first result I would have edited correct copy to satisfy a broken test. Files: `research/Message_Stacyann_Young_2026-08-27.md` and a 2-page PDF.

2026-08-27 — TEN CONTRIBUTOR-LINK REMINDER EMAILS GENERATED, AND WRITING THEM EXPOSED A REAL HOLE IN THE WITHDRAWAL SCANNER

`scripts/build_reminder_emails.py` writes one file per person for **E-10, E-14, RR-106, RR-109, RR-110, RR-116, V-AI-08, V-AI-12, V-AI-23, V-AI-27**. **Nothing is hand-copied**: names, codes and unguessable links come from `research/Contributor_Links.md`, which is itself generated from `api/_contributor-roster.js`, and **the fallback date is read out of `api/contributor.js` at `FALLBACK_DATE`**, which returns *Saturday, 5 September 2026*. A date typed ten times is a date that drifts, and a reminder carrying a wrong link is worse than no reminder. Each email is ~150 words, thanks the person for **the study they actually did** (reliability ratings, comparison study, or detection panel, chosen by code prefix rather than one form letter), states the link takes about a minute, names the fallback date, and says plainly that **choosing to stay anonymous is a perfectly good answer and the form handles it in one click**. **THE HOLE**. The first run wrote V-AI-08's file with **her full name in the path and her first name in the greeting**. `withdraw_contributor.py --check` reported **clean**, and it was wrong twice over. **The register holds four name forms for V-AI-08: full name, short-form full name, surname and short first name**. The short first name is a nickname, and its word-boundary pattern does not match the longer first name it abbreviates, so the bare first name passed untouched. **And the scan read file CONTENTS only, so a withdrawn name sitting in a FILENAME was invisible entirely**. A name is exposed by a directory listing just as surely as by a paragraph. **The filename scan is now part of `scan_traces` and is demonstrated**: reinstating the badly named file makes the check report that path at line 0 with the marker "*(in the filename)*". **MY FIRST FIX WAS WRONG AND THE FILE HAD ALREADY WARNED ME**. I

added the bare long-form first name to the register. It produced **77 traces**, because that form collides with **a different person entirely, an active collaborator carrying 62 mentions across the tree**, against 3 for the withdrawn contributor. `scripts/withdraw_contributor.py` already carried a comment warning about exactly that collision and naming the collaborator, and I overrode it before checking. Reverted within the hour with the reasoning recorded in the register so the next person does not repeat it. **THE GREETING PROBLEM IS NOW SOLVED WHERE IT BELONGS, IN THE GENERATOR.** V-AI-08's reminder is written as `V-AI-08.md` with no name in the filename and a neutral *"Hi there,"* greeting. Her reinstatement of 2026-08-19 is scoped to the contributor link and to four named files; **a reminder about that link is within the reinstatement, but a NEW file carrying her name is not**, and the file does not need to identify her to do its job since the owner sends it to an address he already holds. Suite 92 checks, register check exit 0.

2026-08-28 — THE REMINDER-EMAIL COMMIT WAS BLOCKED BY MY OWN TWO GUARDS AND BOTH CATCHES WERE CORRECT

The ten contributor-link reminder emails were finished and the commit was refused by the pre-commit hook with exactly two FAIL lines, **92 checks, 2 failed, 2 skipped. NEITHER WAS A FALSE ALARM, WHICH IS THE POINT WORTH RECORDING. (1) no withdrawn contributor name survives fired on research/MASTER_TRACKER.md:2594**. My own log entry describing the withdrawal-scanner fix had spelled the withdrawn contributor's name out four times while explaining why her name must not appear in generated files. The register's rule for historical surfaces is that **the log keeps the decision and loses the name**, and the tracker is not one of the four files her name is permitted in. The entry is rewritten to participant code and structural description only: the four register name forms are described by shape rather than quoted, the offending filename is referred to by its marker (`in the filename`) rather than reproduced, and the 62-mention collision is described as a different, active collaborator with the reader pointed at `scripts/withdraw_contributor.py`, which is the one place the names legitimately live. **Nothing about the lesson was lost; only the name was.** Register check now exits 0 with *"No trace of any withdrawn contributor remains outside the register."* **THIS IS THE SECOND TIME IN ONE DAY THE SAME GUARD CAUGHT ME DOING THE SAME THING**, which is a fair measure of how easily a name leaks back in through the very document that records its removal. **(2) Master Tracker written to today fired with "NO ENTRY for 2026-08-28; newest is 2026-08-27"**. The date rolled over mid-task. **That guard was written yesterday, in this session, in direct answer to Phillip asking how he could trust that the tracker was being kept**, and its first live catch is on my own commit rather than on a hypothetical one. That is the behaviour he was owed: the mechanism did not depend on my remembering. This entry is what clears it. **DELIVERED THIS TURN:** `scripts/build_reminder_emails.py` and the ten files in `research/Reminder_Emails_2026-08-27/`, nine named for their recipient and **V-AI-08's written as**

V-AI-08.md with a neutral greeting, sent to Phillip as files rather than as an artifact, since artifacts render in a sandbox that blocks downloads.

2026-08-28 — THE STUDY LABEL IN THE REMINDERS WAS INFERRED FROM A CODE PREFIX, AND VALIDATING IT FOUND A REAL BUG

Each reminder thanks the person for the study they actually did, and that label was chosen by a hand-written prefix map: E to the reliability study, RR to the comparison study, V-AI to the detection panel. **That is the same hand-built-list defect that produced the phantom Rungs paper and hid the accepted CEP article yesterday**, so it was checked rather than trusted. **NEITHER OBVIOUS SOURCE SETTLES IT.** `api/_contributor-roster.js` carries a declared `kind`, but `kind` is `panel` for **BOTH the comparison study and the detection panel**, so it holds less information than the prefix does. `research/Contributor_Links.md` carries a free-text note that does name the study, but the note is **blank for 15 rows, including 3 of the 10 people being reminded**, so it cannot be read off per person either. **THE MAP IS THEREFORE VALIDATED INSTEAD OF ASSUMED:** every roster row that does carry a note must agree with it or the build exits non-zero. Result: **24 notes agree, 0 disagree, 15 blank rows take the validated label.** A map that agrees with 24 declared rows is evidence for the blanks; a map nobody checks is a guess, and a guess inside a thank-you is worse than no thank-you. **THE VALIDATOR FIRED ON ITS FIRST RUN AND FOUND A LATENT BUG I HAD NOT LOOKED FOR.** `group_of` returned `E` for anything that was not V-AI or RR, so it classified **E-08, M-01 and V-HR-01, who are co-authors and facilitators**, as reliability-study expert raters. None of the three is on the reminder list, so no wrong email was ever written, but the function was unsound and would have thanked a co-author for ratings she never gave. **Fixed by declaration rather than by luck:** participation is now gated on `kind` in `{panel, rater}`, authors and facilitators are excluded by the roster's own field, and the prefix is used only to split the two panel arms. The build also refuses outright if any target code is not a study participant. **COMPLETION VERIFIED BEFORE THANKING ANYONE, per CLAUDE.md VIII:** `research/check_completion.py` returns **COMPLETE for all eight panel codes**, V-AI-08, V-AI-12, V-AI-23, V-AI-27, RR-106, RR-109, RR-110 and RR-116, read from the anon-readable aggregate views. E-10 and E-14 are reliability-study raters on a different instrument and the checker correctly reports *"NO ROW"* for them rather than passing them silently. **A SECOND SMALL DEFECT FIXED:** the subject line read *"Quick one, there*

- *your name on the JRS write-up"* for the withheld-name recipient, which is not a greeting

anyone writes. The subject now drops the name entirely when the name is withheld. **LINKS AND DATE PROVEN, NOT ASSUMED:** all **10 keys checked back against** `api/_contributor-roster.js` and **all 10 match, all 10 distinct**, and the fallback date *Saturday, 5 September 2026* is present in all 10 and is read

from `FALLBACK_DATE` in `api/contributor.js` rather than typed. **MY FIRST VERIFICATION PROBE WAS BROKEN AND I DID NOT ACT ON IT.** A shell one-liner with a malformed awk field separator reported **0 of 10 links matching**. The separator was the fault, not the data. Rewritten against the roster source in Python: 10 of 10. **That is the fourth broken probe in two days, and in every case trusting the first red result would have meant editing correct output to satisfy a defective test.** Suite 96 checks, 0 failed. Register check exit 0. Commit `fc932b` carried the emails; this turn hardens the generator.

2026-08-28 — THE CORRECTED SKIP-CI HOOK CONFIRMED ON THREE MORE COMMITS, INCLUDING THE LONGEST MESSAGE YET TESTED

PR #10 returned bot status for all three of today's pushes and **Cloudflare reported Deployment skipped on every one:** `fc932b`, `b7c092f` and `1c95d3b`. **Vercel reached Ready on each**, which again confirms the token costs nothing on the Vercel side. **The new evidence is the message length.** The hook was corrected on 2026-08-26 to insert `[skip ci]` on line 3 rather than append it, and the longest confirmation on record until now was `ea96a85` at 1,758 bytes with the token at byte 77. Today: **`fc932b` at 2,352 bytes, token at byte 65;** `b7c092f` at 2,013 bytes, token at byte 64; `1c95d3b` at 341 bytes, token at byte 62. **The 2,352-byte message is now the longest commit message proven to skip**, and it sits well past the 1,031-byte offset at which `f607e86` and `d07268e` FAILED under the old appending hook. That is the point the fix was meant to establish: **the offset stays fixed near the top however long the body grows**, so the length cap Cloudflare reads under is never reached. `check_skip_token_lands_where_cloudflare_reads_it` asserts the offset stays under 194 and all three commits are comfortably inside it. **Nothing was actioned from these notifications** and none needed it; they are ten routine deployment status updates from `vercel[bot]` and `cloudflare-workers-and-pages[bot]`, recorded here only because they are live evidence for a hook whose behaviour this file documents in detail. **The Cloudflare Git integration is still connected and still a dashboard action to remove:** Workers and Pages, jrsstandardcom, Settings, Build.

2026-08-28 — THE CREDENTIAL SENTENCE WAS VERIFIED AGAINST THE LIVE DATABASE AND THEN PLACED ON THE FOUR PAGES THAT ASK A STRANGER TO ACT AND OFFERED NO PROOF

Phillip supplied the sentence from `access.html` and asked it be verified and integrated where appropriate. **EVERY FIGURE VERIFIES AGAINST `/api/panel-stats` READ LIVE AT 2026-08-28T09:06:23Z:** `reviewers_all` 58 scoped *all three studies*, `completers_all` 36, `countries_all` 16 scoped *all completers*, and the endpoint's own `basis` field states *"completers graded all 24 records in their set"*, which is the 24-record claim. **THE SCOPING IS THE PART THAT MATTERS AND IT IS CORRECT.** The sentence attaches 16 countries to the 36 completers, never to the 58. `geo_note` records that attaching the country figure to the reviewer total is a **recorded past defect**, and `check_panel_geo` exists to prevent it. The sentence as written does not commit it. **IT WAS ALREADY LIVE ON FIVE PAGES AND MISSING FROM THE FOUR THAT NEEDED IT MOST.** Present on `access.html:81`, `investigator-guides.html:110`, `org-pilot.html:154`, `reviewer/index.html:124` and, in a figure-free form, `training.html:3193`. **Absent entirely from `index.html`, `enterprise.html` and `review-engine.html`:** the homepage and both Track 1 commercial pages, which is to say every page where a platform buyer arrives cold and is asked to click, to open a scoping call, or to request a token. **`training.html` carried the credential with NO figures at all,** and no binder, so its enrolment overlay asked for a full name and a work email on authority alone. **`scripts/integrate_credentials.py` places it, and the binder is READ from `access.html` at run time rather than pasted into the script,** so the script cannot become a fifth stale copy of a 3,889-byte block that `check_panel_binder_identical` requires to be byte-identical in what is now 14 places. Dry run by default. **THE FIGURES ARE NEVER TYPED:** every numeral sits in a `data-panel` span the binder overwrites from the endpoint, and the markup numeral is the marked fallback that renders with a dotted underline if the fetch fails. **PLACEMENT FOLLOWS THE DECISION ALREADY RECORDED AT `access.html:78`:** the credential goes BELOW the button row on all four, because it is supporting evidence and not a precondition. **MEASURED ON THE RENDERED PAGE, NOT READ FROM SOURCE:** the block sits at 5.2% depth on `index.html` at 390px, 3.9% on `enterprise.html`, 7.5% on `review-engine.html`, and all three spans report `data-panel-state="live"` rather than stale. No horizontal overflow at 390px or 1280px. **`check.html` WAS DELIBERATELY EXCLUDED AND THE EXCLUSION IS ASSERTED IN CODE.** It already publishes `completers_detection` 16 and `countries_detection` 11, the detection-panel figures. Putting 36 and 16 beside them would place two populations in one viewport, which is precisely the top-versus-bottom mismatch the scoped keys were introduced to end. **`check_trust_pages_carry_their_proof` fails if the all-studies keys ever appear on that page.** **THE NEW GUARD FOUND A REAL DEFECT ON ITS FIRST RUN.** `reviewer/index.html:124` was still bound to the legacy **unscoped keys** `completers` and `countries`. The endpoint returns them as aliases so the page was not visibly broken, but the binder's own

comment states the problem exactly: "a bare `countries` meant two different populations in two paragraphs of the same page, which is the whole of the top-versus-bottom mismatch." **It was the last page in the tree still on them;** migrated to `completers_all` and `countries_all`, and a site-wide scan confirms **zero unscoped bindings remain. The guard is demonstrated rather than asserted:** removing the new block from `index.html` makes it report "*index.html: no credential (the homepage: first contact, and it asks for a click into both tracks)*", restored immediately after. It holds a reason string per page, so a future reader can argue with the list instead of guessing why a page is on it. Suite now **97 checks, 0 failed.** every published panel figure is bound scans **73 HTML files;** panel binder copies are byte-identical now covers **14 pages.**

2026-08-28 — JEFFREY BILLUPS SUBMITTED THE BLIND SECOND READ. THE OWNER COULD NOT TELL WHAT HE DID, AND THE REASON IS A REAL GAP: NO DEPLOYED SURFACE READS THE ANSWERS.

Phillip saw a new row dated 2026-08-28 on the programme status page and asked what it was. **ANSWER:** `activity` and `source` both read `recheck-submit`, timestamped 2026-08-28T04:22:14.765Z, country US, `consent_contact` true, `consent_transfer` false, `consent_public` false. That is `api/recheck.js`, the **blind second-reader instrument for the public-records study**, whose own header states why it exists: "*The manuscript reports 32 reads produced by one person. The one weakness a referee will name is that nobody checked those reads independently.*" `/api/asset-stats` confirms it under `named_professional_engagement.blind_second_read: links_issued 3, submitted 1.` **THIS IS THE ANSWER TO THE QUESTION PUT TO STACYANN YOUNG YESTERDAY.** The message sent to her asked whether she ever got anyone to independently review the determinations, because Section 7 of the FOIL manuscript concedes "*All 32 reads were recorded by a single domain reviewer, so no inter-rater agreement is estimated.*" A second read has now arrived. **WHICH SLOT HE USED IS NOT PROVABLE FROM ANY READABLE SURFACE AND IS NOT ASSERTED.** `research/Blind_Recheck_Links_2026-08-09.md:9-11` records **R1 offered to Stacyann Young on 2026-08-09 to forward to her attorney contact, R2 and R3 unassigned**, which makes R1 the plain reading. But `api/recheck.js` deliberately does not store who holds a key, and `submitted()` in `api/asset-stats.js:94-104` falls back to `created_at` when no slot is present, so **submitted:1 does not by itself prove a slot was recorded.** Stated as inference, not fact. **THE GAP THAT MADE HIM ASK.** `api/recheck.js:163-179` writes the ten answers, the slot, `answered_count`, `prior_familiarity` and `consent_named_in_paper` as JSON into `pilot_contacts.message`. **Nothing deployed reads that column.** `api/people-9dd1ecdf6f8cdfd4` returns the row with `detail:"";` `api/asset-stats` returns a count; `api/leads-4b7e2c9af106d385` correctly excludes

it as non-commercial. **The single most valuable research event of the month landed where the owner can see that it happened and not what it said.** `scripts/score_blind_recheck.py` closes it. Pulls every `source='recheck-submit'` row, scores each against `research/Blind_Recheck_KEY_E08.md` (never deployed), and reports per-case agreement, percent agreement and **Cohen's kappa. Kappa and not raw agreement alone, because the key is 6 Ready, 3 Needs work, 1 Gap: a reader who answered Ready ten times scores 60% and has demonstrated nothing.** Verified on synthetic input without touching the database: perfect 10/10 gives 100% and kappa 1.0; **all-Ready gives 60% and kappa 0.0**; a partial 4-of-10 return gives 75% and kappa 0.636. Vocabulary, project URL and label set are all read out of `api/recheck.js` rather than restated. **FAIL-CLOSED, NOT GUESSED.** `pilot_contacts` has RLS on with no anon read, and no service key exists in this environment, so the script exits 1 with `[REQUIRED_ENV_PARAM]` naming the three accepted variable names and stating the key lives in the Vercel environment and must not be committed. **The ten answers remain unread; nothing about their content is claimed. A SECOND, UNRELATED DEFECT WAS FOUND WHILE READING THAT ROW AND IS FIXED.** `api/people-9dd1ecdf6f8cdfd4.js:193` set `training_completed_on` to the row's own `created_at` for every non-enrolment row, so **46 of 58 rows carried a training completion date while `training_completed` was false** and every one of those dates equalled the row date. The owner table at `programme-status-9872fb93cc94.html:1451` guards the field on `training_completed` and therefore looked correct. **The CSV export at line 1515 does not guard it**, so a downloaded file asserted 46 completions that never happened. Fixed at source so both surfaces are right. `training_completed_named_count` was never affected: it filters on `training_completed`. `check_completion_date_implies_completion` demonstrated to FAIL with *"training_completed_on falls back to r.created_at with no completion test"* against the pre-fix expression. Suite now **98 checks, 0 failed. ONE ANOMALY NAMED AND NOT EXPLAINED:** `blind_second_read` reports **links_opened 0 with submitted 1.** `api/recheck.js:112-127` writes an open ping on GET unless `?src=owner|verify|test|selftest|deploytest`, and the open window opened 2026-08-09, well before this submission. A submitted packet that was never recorded as opened means either the ping was suppressed by a src tag on the link he followed, or the write failed inside the try/catch that is designed never to block the packet. **Not resolved from a readable surface, so it is logged rather than explained.**

2026-08-28 — DID THE SECOND READER FINISH? THE HONEST ANSWER WAS THAT NOTHING DEPLOYED COULD SAY, AND THAT IS NOW FIXED RATHER THAN ANSWERED BY HAND

Phillip asked whether Jeffrey Billups completed the assigned task and provided results. **WHAT IS PROVEN: HE SUBMITTED. WHAT WAS NOT PROVABLE: WHETHER HE FINISHED.**

`api/recheck.js:150-153` accepts a partial return on purpose, with the reason written into the file: *"Unanswered cases are accepted rather than rejected: a partial return is data, and forcing ten before anything can be saved risks losing all ten."* So a `recheck-submit` row proves arrival and says nothing about completion. The `answered_count` and the ten labels sit in `pilot_contacts.message`. **RLS CONFIRMED BY PROBE, NOT ASSUMED:** the public anon key from `research/check_completion.py:35` returns **HTTP 200 with []** on `pilot_contacts?source=eq.recheck-submit`, and `recheck_progress`, `recheck_results` and `recheck_agreement` all return **PGRST205, no such table**. There is no aggregate view for this instrument. **HE IS IN NO ASSIGNMENT RECORD.** A corpus-wide search for the name returns exactly one hit, my own tracker entry from earlier today. `research/Blind_Recheck_Links_2026-08-09.md:9-11` shows **R1 offered to Stacyann Young to forward to her attorney contact, R2 and R3 unassigned**, and the forwarded recipient was never named here, which is consistent with an inbound submission from someone the repository has never held. **THE FIX IS NOT TO ANSWER THE QUESTION, IT IS TO MAKE THE DASHBOARD ANSWER IT.** `api/asset-stats.js` already parses `pilot_contacts` server-side, so `blind_second_read` now publishes **complete_returns**, **partial_returns**, **answers_recorded**, **cases_offered** and **unparsed_rows** beside `submitted`. Completeness is now readable without a service role key on his laptop. **AND IT PUBLISHES COUNTS ONLY, WHICH IS THE HARD CONSTRAINT.** No label, no case, no agreement figure and no kappa goes into that endpoint. **An agreement percentage sitting beside a public ten-case list reconstructs the answer key**, and the blind is the entire instrument. `check_second_read_completeness_is_published` asserts both halves: that `complete_returns` is present, and that no field name matching agreement, kappa, label, per_case, score, correct or key ever appears. Demonstrated both ways: stripping `complete_returns` reports *"does not publish complete_returns"*, and adding `agreement_pct: 80` reports *"publishes the field 'agreement_pct', which leaks the answer key"*. **MY FIRST VERSION OF THAT GUARD WAS A BROKEN PROBE AND I CAUGHT IT BEFORE SHIPPING.** It scanned the whole block for the substring *agreement* and failed on the note I had written explaining why no agreement figure is published. A guard that fires on its own documentation would have had me delete a correct explanation to satisfy a bad test. Rewritten to parse **field names only**. That is the fifth broken probe in three days and the fifth time the first red result was wrong. **A SECOND DEFECT FOUND IN THE SAME BLOCK AND FIXED.** The suppressed-cohort entry hardcoded `sent: 0` with the reason *"Awaiting the second reader being named. None has been sent."* while `submitted: 1` sat in the same object. **A returned packet is proof a link reached a reader.** State is now derived: a submission moves the cohort from SUPPRESSED to ACTIVE by itself, and `sent` is **null rather than 0**, because the links are forwarded by hand and this system never observes the send. A zero asserted a fact; null states the truth, which is that the send is unobserved. Suite now **99 checks, 0 failed**. `scripts/score_blind_recheck.py` still holds the per-case detail and still fail-closes without the service key; **the ten answers remain unread and nothing about their content is claimed.**

2026-08-28 — THE FOIL PAPER WAS CITING THE COMPANION STUDY'S EXCLUDED-CASES SENSITIVITY ANALYSIS AS ITS HEADLINE CROSS-DOMAIN RESULT. FIXED IN TWO PLACES, PLUS A DELIVERY DEFECT IN YESTERDAY'S PDF.

Phillip supplied the CFOC submission and the revised FOIL PDF and asked for the research to be completed and the article revised. **THE UPLOADED CFOC DOCX IS THE REPOSITORY'S OWN `research/CFOC_Submission_2026-08-08.md` EXPORTED TO WORD**, verified sentence by sentence: the only difference is a dropped `---` separator. A duplicate I had extracted was deleted rather than kept. **THE REAL FINDING IS A FIGURE THAT WENT SUPERSEDED IN THREE DOCUMENTS AT ONCE.** `FOIL_Article_Draft.md` section 5.6 and its findings summary at line 29 both cited the employment corpus at "22 cases from 22 distinct sources", "7 of 9 against 2 of 13, $p = 0.0073$, odds ratio 19.25" and "6 of 8 against 1 of 8, $p = 0.041$, odds ratio 21.0". Every one of those is computed on the 22-case SCREENED set. **The employment corpus was corrected on 2026-08-24:** two matters fail the stated inclusion criteria and the analysis runs on 20 (`Employment_Records_Article_ISACA_2026-08-21.md`, notes 2 and 5), where the primary association is $p = 0.0194$, odds ratio 15.00, 6 of 8 against 2 of 12. **THIS IS WORSE THAN A STALE NUMBER.** That manuscript states outright: *"Including them produces $p = 0.0073$ with an odds ratio of 19.25. Because those matters do not meet the stated inclusion criteria, this result is reported only as a sensitivity analysis."* **The public-records paper was publishing the companion study's sensitivity analysis as its primary cross-domain evidence. AND ONE OF THE TWO EXCLUDED MATTERS IS A PUBLIC-RECORDS ADVISORY OPINION**, appendix A15, FOIL-AO-19774, excluded precisely because it belongs to the corpus this paper reports. A referee opening the companion manuscript would have found a public-records case propping up the public-records paper's cross-domain claim. That is the most damaging form the error could take. **PROVENANCE WAS ESTABLISHED BEFORE ANYTHING WAS REWRITTEN, FROM LIVE DATA.** `scripts/recompute_sustained_coding.py` pulls the 22 screened employment matters from `bench_outcomes` and recomputes the sustained coding with Fisher's exact written out by hand, because scipy is not installed and a p value quoted to a federal council must not depend on a package being present: **6 of 8 against 1 of 8, $p = 0.0406$, odds ratio 21.00**, reproducing the quoted figures exactly. **The numbers were never wrong; their basis was superseded. THE 20-CASE SUSTAINED CODING WAS NOT COMPUTED AND THE SCRIPT SAYS SO RATHER THAN GUESSING.** Its exclusion screen flagged 22 rows where the manuscript names 2, because the appendix-A-to-row mapping is not in the anon-readable data. It exits 2 with `[REQUIRED_ENV_PARAM]` and refuses to drop two rows by inference. **The published $p = 0.0291$ on 13 resolved matters is cited from the manuscript instead. THREE DOCUMENTS CORRECTED:** both occurrences in `FOIL_Article_Draft.md` (4,344 to 4,369 words), and the CFOC outreach paragraph, which had gone out under Stacyann Young's name to the Chief FOIA Officers Council and a named DOI attorney carrying 22 adjudicated cases, **six of eight against one of eight, $p = 0.041$.**

PROVEN NOT TO HAVE DISTURBED THE PAPER'S OWN FIGURES: every reported figure token was extracted before and after and diffed. **8 removed, all employment; 6 added, all their corrected counterparts; 13 unchanged, all public-records.** All 9 of Stacy's edits intact. **MY FIRST REPLACEMENT PROSE WAS WRONG AND I CAUGHT IT.** It read *"the figures previously cited here ... are superseded"*, which is a note to an editor, not manuscript prose: a referee has no idea what was previously cited. Rewritten as clean text; the superseded figures live in the commit and here. **A SEPARATE AND SERIOUS DELIVERY DEFECT.** `scripts/render_report_pdf.py` wraps its source in `<body>` and **does no markdown conversion at all.** Handed a `.md` manuscript it produces a PDF with literal `#`, `and` `---` markers and every heading, table and paragraph collapsed into one wall of running text. The 11-page FOIL PDF delivered to Phillip on 2026-08-27 has that defect; **it was reported as a manuscript and it was unformatted source.** Caught here only because a longer document rendered to 6 pages instead of 11, **which did not add up, and the PDF was rendered back through the browser and read rather than trusted.** `scripts/md_to_html.py` supplies the missing step, with no external dependency because markdown, mistune and commonmark are all absent here: ATX headings, bold, italic, inline code, superscript already written as HTML, pipe tables, ordered and unordered lists, rules, block quotes and paragraphs. **Output on this manuscript:** 1 h1, 11 h2, 12 h3, 4 tables, 78 paragraphs, 0 unconverted markers. **The renderer now routes any `.md` through it, so no caller can forget.** Re-rendered: 99,284 bytes, 11 pages, verified by screenshot to carry real headings and a title block. Two guards added and both demonstrated against the pre-fix state: `check_crossdomain_citation_is_current` holds a map of seven superseded fragments to the reason each is wrong and fails with *"still cites '22 cases from 22'"* and three more; `check_markdown_pdfs_are_converted` fails with *"does not route a .md source through md_to_html.py"*. `scripts/audit_cfoc_claims.py` verifies all 12 empirical claims in the outreach emails against the manuscripts and now passes. **Two of its own rules were false positives I fixed rather than acted on:** it searched the paper for the email's phrasing *"no relationship"* when the paper writes *"is null (p = 1.000)"*, and a bare `/certif/` fired on Stacyann Young's genuine SUNY and New York State Archives certifications. That is the sixth broken probe in three days. **Artifacts:** `FOIL_Paper_REVISED_2026-08-28.pdf` 11pp, `FOIL_Article_REVISED_2026-08-28.docx`. Suite now 101 checks, 0 failed. STILL OPEN**: the blind second read is complete at 10 of 10 but unscored, so Section 7's single-reader limitation stands unchanged and correctly.

2026-08-28 — THE SECOND READ IS SCORED AND THE ARTICLE IS UPDATED THROUGHOUT. 70.0 PERCENT AGREEMENT, COHEN'S KAPPA 0.474. THE SINGLE-READER LIMITATION IS NARROWED, NOT RETIRED.

Phillip required the article updated with the re-examination. The blocker was real and was removed rather than argued with: the ten answers sat in `pilot_contacts.message` behind RLS, and **the service role key exists in the Vercel environment even though it does not exist here.** `api/recheck-answers-b1a768e88d3e48bd.js` is a third owner-only endpoint on the established pattern, opaque slug, no token, no analytics tag, never linked. It returns the reader's labels, reasons, slot and consent flags and **deliberately does not contain the answer key**: the original reads stay in `research/Blind_Recheck_KEY_E08.md`, which is never deployed, so a leak of this slug exposes one person's labels and still leaves nothing to score them against, and the blind on the two unissued packets survives. Deployed to `main` at `ac43692`. **THE READER IS JEFFREY BILLUPS, SLOT R1**, which is the packet offered to Stacyann Young on 2026-08-09 to forward to her attorney contact, submitted 2026-08-28T04:22:14Z, **10 of 10 answered**, prior familiarity with the instrument recorded as *"None / Independent reviewer"*, and **he reported knowing the documented outcome in 0 of the 10 cases**, which is the blind holding. **He consented to be named in the paper.** **THE RESULT, COMPUTED AND NOT ESTIMATED:** exact agreement **7 of 10, 70.0 percent, 95 percent Wilson 39.7 to 89.2; Cohen's kappa 0.474 unweighted**; linear weighted kappa 0.559; Gwet's AC1 0.582. **ALL THREE DISAGREEMENTS WERE ADJACENT AND NONE WAS A READY AGAINST A GAP.** Case 1 Ready to Needs work, case 4 Ready to Needs work, case 5 Needs work to Ready: **the second reader was stricter on two and more lenient on one, and every disagreement sits on the Ready and Needs work boundary**, which is the boundary the instrument is least sharp about. **The single Gap read, which carries the operational consequence, was reproduced exactly.** **THREE COEFFICIENTS ARE REPORTED AND THE UNWEIGHTED KAPPA LEADS, WHICH IS THE LOWEST OF THE THREE.** Reporting only AC1 at 0.582 would be choosing a statistic after seeing the data. The weighted kappa is justified because the scale is **ordinal** and every disagreement was adjacent; AC1 is justified because **Ready holds 6 of 10 of the margin**, the condition under which kappa is known to understate. Each is stated with its n. **THE LIMITATION NARROWS, IT DOES NOT VANISH, AND THAT IS ENFORCED IN CODE.** Section 7 now reads that **10 of 32, not all 32, were re-read**, that 0.474 is moderate and is *"evidence that the read is not idiosyncratic to one person, and not evidence that two readers would classify the full corpus alike"*, that the interval is wide because ten cases cannot make it narrow, that **a single re-read cannot separate reader dependence from case difficulty**, that the remaining 22 cases carry the original limitation in full, and that **two further packets were prepared and have not been returned.** **FIVE SECTIONS CHANGED, EVERY FIGURE READ FROM JSON AND NONE TYPED:** Abstract, Data availability, Methods 4.6 (new, describing what the reader was and was not shown), Results 5.7 (new), Limitations. `scripts/apply_second_read_to_manuscript.py` refuses to run if `research/Blind_Recheck_RESULT_2026-08-28.json` is absent. **4,369 to 5,045 words. PROVEN**

NOT TO HAVE DISTURBED ANYTHING ELSE: the figure-token diff against HEAD shows **0 removed** and only the seven new agreement figures added. **All 9 of Stacy's edits intact.** `check_second_read_reported_honestly` asserts all five computed figures are present, that **the lowest coefficient is reported**, and that Limitations still says "*not all 32*" and "*is not a panel*". Demonstrated both ways: replacing 0.474 with 0.582 fails with "*the lowest of the three coefficients (0.474) is not reported*", and softening "*is not a panel*" fails with "*a subset re-read is being presented as if it settled the corpus*". **The guard raised a NameError on first run and was fixed rather than deleted.** Artifacts: `FOIL_Paper_REVISIED_2026-08-28.pdf` **13 pages**, `FOIL_Article_REVISIED_2026-08-28.docx`, `Blind_Recheck_RESULT_2026-08-28.json`. Suite now **102 checks, 0 failed**.

2026-08-28 — ALL 18 SURGICAL CORRECTIONS APPLIED, AND VERIFYING ITEM 11 EXPOSED A REAL ANALYTIC DEFECT THAT NEITHER OF THE OWNER'S TWO OPTIONS COVERED

Items 1 to 9 are KEEP instructions and are asserted rather than written: `scripts/apply_final_surgical_list.py` fails if any of the nine co-author approved strings is missing after the pass. **9 of 9 verified present. ITEM 11 WAS LEFT OPEN ON PURPOSE AND THE ANSWER IS A THIRD THING.** He wrote that the 27-case wording should be used "*only if that accurately reflects the actual analytic design. If the actual exclusion was based on missing notes rather than document class, the sentence must reflect that instead.*" **Live bench_outcomes for the public-records corpus: Ready 18 cases / 17 noted, Needs work 9 / 7, Gap 5 / 4. 32 cases, 28 notes, 4 without.** The Section 5.3 table was drawn on $n = 9$ and $n = 18$, all 27 case-level sources, but only 24 of those 27 carry a note, so **THREE CASES WITH NO NOTE WERE SITTING IN THE "NOT STATED" COLUMN.** Absence of a note is not a note that fails to state a reconstructability failure, and coding it as one inflates the comparison group. **Both restrictions are now stated: case-level first, then note-carrying. THE CORRECTION MAKES THE RESULT STRONGER, WHICH IS WHY IT HAD TO BE CHECKED RATHER THAN ASSUMED HARMFUL.** As published, `[[6,3],[0,18]]` on 27 gives $p = 0.00028$, which is arithmetically right for the table as drawn. Restricted to the 24 noted cases, `[[6,1],[0,17]]` gives $p = 0.0000520$. Both recomputed with Fisher's exact written out by hand, `scipy` being absent. Cell counts are forced arithmetic, not a re-reading: a coded note must exist, so the 6 stated Needs work cases all carry notes. **TWO PLACES THE OWNER'S LIST DID NOT REACH AND THAT WOULD HAVE LEFT THE PAPER CONTRADICTING ITSELF. (1) The ABSTRACT carried the same table, "six of nine ... against none of eighteen ... $p = 0.00028$ ".** Correcting 5.3 and leaving the abstract would have put the paper's two most-read passages in conflict, which is the defect item 10 exists to remove. Same correction, same source. **(2) ITEM 9 WAS VIOLATED IN EXACTLY ONE PLACE AND IT WAS NOT THE OBVIOUS**

ONE: the abstract described the employment corpus as *"flagged records"* against *"passed records"*. That corpus is read with the same five-condition instrument, so those are JRS classifications and now read *"Needs work or Gap"* and *"Ready"*. **Line 17's "can read as complete" was left alone:** ordinary prose, not a classification, the same judgment the co-author's own terminology pass made. **ITEM 10'S OPENING SENTENCE NEEDED THE OWNER'S NUMBER, NOT MY COMPUTED ONE.** My first pass wrote *"For the 24 cases with contemporaneous basis notes"*, which contradicts Section 5.1's 28 in precisely the way item 10 forbids. **28 is corpus-wide, 24 is the coded subset.** The opening now uses 28 as he specified and the restriction paragraph explains the drop to 24, so both numbers appear in order and neither surprises. **REMAINING ITEMS APPLIED AS SPECIFIED:** 12 adjudicator to **independent government auditor** in RQ2, 4.5 and 5.2; 13 **Four to Five analyses** now that 5.7 exists; 14 Section 7 to **20 adjudicated with 22 screened**; 15 *"That answers"* to preliminary evidence; 16 *"establishes three things"* to *"provides evidence for three propositions"*; 17 **applied although optional**, abstract softened to *"preliminary evidence that the read responds to the reconstructability property it is designed to assess"*; 18 the unsupported causal implication removed from the introduction. **5,045 to 5,189 words. BOTH COUNCIL EMAILS UPDATED TO MATCH:** the construct sentence corrected the same way, the blind second read added as a **fourth finding** with agreement, both kappas and the adjacency of every disagreement, and *"Three findings"* corrected to four. `scripts/audit_cfoc_claims.py` now verifies **15 claims across both manuscripts, all OK, vocabulary clean**; its stale rules were rewritten to the corrected figures rather than left passing on absence. **MESSAGE TO STACYANN YOUNG, 762 words**, carrying the owner's three instructions verbatim in substance: her friend completed the work, the article and emails are updated, **she should send both emails solo** because they are stronger from twenty years in New York City government and her certifications than from a joint signature, and **he handles the journal submission once she gives final approval**. It reports the agreement honestly including that it does **not** retire the limitation, explains why the construct correction strengthens the result, and raises the two unissued packets as entirely her call. Artifacts: `FOIL_Paper_FINAL_2026-08-28.pdf` 13pp, `FOIL_Article_FINAL_2026-08-28.docx`, `CFOC_Submission_2026-08-08.docx`, `Message_Stacyann_Young_2026-08-28.docx`. Suite **102 checks, 0 failed**.

2026-08-28 — THE MANUSCRIPT'S REPRODUCIBILITY CLAIM WAS FALSE WHEN CHECKED, AND THAT IS THE FINDING OF THIS PASS. FIXED, PLUS A FULL SUBMISSION-CONTROL PACKAGE.

Phillip asked for a final check of References and Data availability against the actual source files before emailing the editor, on the ground that the paper makes unusually strong reproducibility claims. **It did, and one of them did not hold. THE DATA AVAILABILITY STATEMENT CLAIMED:** *"Every figure in*

Section 5 is reproduced by an analysis script using only the Python standard library. **THAT WAS FALSE ON TWO COUNTS.** `research/analysis_foil_2026-08-08.py` exists and runs, but it **covers R1 to R4 only**, so Sections 5.6 and 5.7 were not reproduced at all, and its R2 still computes the superseded construct table, 6 of 9 against 0 of 18 at $p = 0.00028$, which the manuscript corrected earlier today to the 24 note-carrying case-level sources at $p = 0.0000520$. **A reproducibility claim the supporting script does not substantiate is the worst defect a paper about traceability can carry**, and it would have been trivial for a referee to test. `research/analysis_foil_2026-08-28.py` **closes it**: covers 5.2 through 5.7, standard library only with Fisher's exact, the Wilson interval, Cohen's kappa and Gwet's AC1 all written out rather than imported, and **it verifies every figure against the manuscript text on each run**. `--verify` exits non-zero on any mismatch. Current state: **19 probes, 0 mismatches**. The Data availability statement now names that file and `Blind_Recheck_RESULT_2026-08-28.json` explicitly. **MY 5.4 GROUPING WAS THE ONE THING I GOT WRONG AND THE VERIFIER CAUGHT IT.** A keyword screen over the note text produced a degenerate table at $p = 1.00000$ against the manuscript's 0.00466. The grouping is not reliably derivable from the note; it is **declared** in the 2026-08-08 script as Group A 6 Ready / 1 not against Group B 0 Ready / 7 not. Carried forward verbatim with the source cited. **Inference replaced with declared data, for the seventh time in four days.** **MANUSCRIPT ITEM 2 APPLIED:** Section 5.2 "*What it establishes*" to "*What it demonstrates*". Items 1, 3, 4, 5, 6 verified unchanged as instructed. **REFERENCES AUDIT CLEAN AND IT RECONCILES TO 32:** 18 New York appellate and trial decisions, 7 Committee on Open Government advisory opinions, 2 Connecticut FOI Commission decisions, **which is exactly the 27 case-level sources**, plus 5 compliance audits. `scripts/audit_references_and_availability.py` also confirms every count quoted in Data availability matches the live database and **fails outright if the statement ever names the superseded script again.** **SUBMISSION PACKAGE BUILT, `research/JCI_SUBMISSION_2026-08-28/`, 12 files in the owner's six-folder structure**, zipped at 150,207 bytes. **CSV rather than XLSX**, because openpyxl is not installed here and a submission dataset must not depend on a package the next machine may lack. **EVERY FIELD IS LABELLED BY PROVENANCE:** DATABASE for read, note, outcome, citation, URL and collection date; DERIVED for the inclusion flags, computed from the analysis rules rather than hand-assigned; DECLARED for the Section 5.3 coding; and **[NOT IN THE DATASET] for the eight fields the study never recorded**, rather than invented. Jurisdiction, source type, decision date, URL-tested-on, verified-by, and the second-read join are all marked absent. **A package about traceability must not fabricate the one thing it exists to prove. THE BUILDER ITSELF LOST A CODED CASE AND I CAUGHT IT BEFORE SHIPPING.** The first run reported 5 coded Yes where the manuscript says 6: `FIC2012-276` appears only inside its URL, and I was matching against the citation half of the source string. Fixed to match the whole string. **That failure is precisely the class of defect the package exists to detect, and it was in my own builder.** **IT ALSO RESOLVED AN AMBIGUITY I FLAGGED EARLIER AND COULD NOT SETTLE:** of the three Needs work cases not coded Yes, `FIC2015-122` is the one that carries a note and does not state a failure, while `F0IL A0 19646` and `2025 NY Slip Op 00220` carry no note at all. That is now recorded rather than guessed. `scripts/presubmission_audit.py` runs the owner's section XIV checklist as executable assertions: **26 checks, 0 failed**, against the live database, the manuscript text and the built package. It covers all 32 URLs, the 18/9/5 and 15/7/5/5 reconciliations, the 28 notes, the 24 coded, the 7 + 17 split, 6 of 7 and 0

of 17, every p value, the second read's four coefficients, the adjacency of all three disagreements, the disclosure, and that the title is unchanged. **NOTHING NEW WAS MEASURED.** No case re-read, no note re-coded, no reader re-contacted. Suite **102 checks, 0 failed.**

2026-08-28 — THE PACKAGE REACHED OUTSIDE ITSELF, WHICH IS THE ONE FAILURE THIS PAPER CANNOT SHIP. NOW SELF-CONTAINED, AND URL TESTING FOUND A BROKEN CITATION.

The review of the delivered ZIP was correct on every RED item. **The reproduction script queried a live database over the network, embedded an API key, and verified against research/FOIL_Article_Draft.md, a path that does not exist inside the package.** A reviewer on a clean machine could not run it. **ALL SEVEN RED ITEMS CLOSED, AND THE OFFLINE CLAIM IS PROVEN RATHER THAN ASSERTED.** 04_REPRODUCTION/analysis.py reads only files inside the package. Run from the **unzipped copy with socket.socket, create_connection and getaddrinfo all replaced by raising stubs and every proxy variable unset: 20 probes, 0 mismatches, exit 0. NOTHING HARD-CODED THAT THE DATA CAN PRODUCE.** Section 5.3's cells are computed from JCI_JRS_Construct_Coding_Frame.csv and Section 5.4's groups from a **new JCI_JRS_Structural_Coding_Frame.csv**, so the chain is case, coding, analysis, result. The constants nw_stated, rd_stated = 6, 0 and GROUP_A_READY = 6, 1 are gone. **NO CREDENTIAL TRAVELS WITH THE SUBMISSION.** Verified on the unzipped ZIP: **0 files containing sb_publishable, supabase or apikey; 0 external research/ paths; 0 references to the superseded script; 0 occurrences of [NOT IN THE DATASET].** URL TESTING FOUND A REAL BROKEN CITATION, WHICH IS WHY THE INDEX EXISTS. PR-28's stored URL is truncated by one character, compliance-freedom-information-law-requirement, and returns **HTTP 404**; the plural form returns **HTTP 200**. Corrected in the index with the evidence recorded. **AND NINE 403s WERE NOT RECORDED AS BROKEN, BECAUSE THEY ARE NOT.** All nine are nycourts.gov and law.justia.com. **Retried with a full browser user agent and still 403**, so this is host-level refusal of automated requests, not a dead link. Recording them as inaccessible would have been false and would have understated the corpus. They read *"Yes to a person; this host refuses automated requests"*. **32 of 32 verified. THE BLIND SECOND READ JOINS TO THE CORPUS EXACTLY.** The ten packet UUIDs in the never-deployed answer key match **bench_outcomes.id 10 of 10** and **record_id 0 of 10**, so the master dataset's blind-review columns are now populated from real data rather than dropped. **BOTH EMPLOYMENT EXCLUSIONS ARE NAMED**, which I could not do this morning: **F0IL-A0-19774** and the unidentifiable Employment Tribunal entry both match a row by exact citation. The companion file now shows **22 screened, 20 included, 2 excluded** with the reason on each, plus tested URLs. **ADDED:**

00_MANIFEST.txt, 02_DATA/JCI_JRS_Data_Dictionary.txt defining every column in every CSV, 01_MANUSCRIPT/manuscript_verification.txt as the local verification target, and a rewritten README.txt that describes the actual ZIP. 16 files, was 12. MY OWN OFFLINE TEST FAILED FIRST AND THE SCRIPT WAS FINE. An exec() harness broke on __file__; re-run properly through a sitecustomize that blocks sockets, it passed. The seventh broken probe in four days, and again the first red result was mine, not the code's. check_submission_package_is_self_contained asserts no credential, no external path, no placeholder and standard library only, demonstrated to FAIL when a credential is planted. Suite 103 checks, 0 failed. presubmission_audit.py still 26 of 26. NOT SENT. The submission remains two attachments, manuscript DOCX and PDF, held for Stacy's final approval; the package is held for an editorial request.

2026-08-28 — THE REVIEW'S TOP RED ITEM WAS RIGHT ABOUT THE SYMPTOM AND WRONG ABOUT THE FIX, AND FOLLOWING IT WOULD HAVE CORRUPTED A CORRECT MANUSCRIPT.

The review reported that Section 5.4's table says 6 of 7 while the packaged structural coding frame produces 5 of 6, and recommended changing the manuscript. The manuscript was right. My coding frame was wrong. ROOT CAUSE, AND THE REVIEW FOUND IT ITSELF WITHOUT CONNECTING IT: PR-10's stored citation reads OIL A0 19746, missing its leading F, so my classifier keyed on "FOIL AO" returned N/A and silently dropped a Ready advisory opinion out of group A. Its own item 5 flagged that typo as unrelated citation hygiene. Restoring it gives group A = 7 cases, 6 Ready, exactly the published table. The typo is corrected with its evidence: the source URL is docsopengovernment.dos.ny.gov/coog/ftext/f19746.htm, the Committee's FOIL advisory-opinion path, matching the f##### pattern of the other six. p = 0.00466 is unchanged, because Fisher on [[6,1],[0,7]] and [[5,1],[0,7]] both return it, which is precisely why the p value could not have caught this. RED 2 AND 3 WERE ARTIFACTS OF MY OWN GENERATOR, NOT MANUSCRIPT ERRORS. The review read manuscript_verification.txt and found "analysisfoil2026-08-28.py" and "BlindRecheckRESULT2026-08-28.json". My markdown stripper removed _ along with and backticks, rewriting every filename in the Data availability statement. Underscores are load-bearing in a filename; the stripper now removes emphasis only. The manuscript's JSON name was correct all along; only the script name was genuinely wrong, because I named the repository file where the package ships analysis.py. Corrected. RED 4 CLOSED: PR-01, PR-03, PR-04, PR-05 and PR-10 were N/A. Four carry NY3d reporter citations, which is the New York Court of Appeals, and PR-01's URL is the /ctapps/ path. The classifier now recognises NY3d, AD3d and /ctapps/. 0 rows remain unclassified. RED 6 AND

7 ARE FLAGGED, NOT FIXED, AND THAT IS DELIBERATE. PR-15 stores 2024 NY Slip Op 0407 against a URL ending 2024_04071, and PR-22 stores 2025 NY Slip Op 0578 against 2025_05783. Appending the digit the URL implies would be inference presented as verification, in a package whose entire purpose is to prevent exactly that, and nycourts.gov refuses automated requests so the decisions could not be read from here. Both are recorded in the source verification index with the implied citation and an explicit *"requires author verification against the published source"*. **Two items the owner must check by hand before sending. AMBER 8 CANNOT BE SATISFIED AND THE FALLBACK IS TAKEN.** The employment corpus holds **zero URLs anywhere in the database**; that study recorded full reporter citations, which are the canonical identifiers for legal sources. The companion file is now explicitly labelled a **citation-based verification record for a separately conducted corpus, not a URL-based reproducibility dataset**, with the reason stated. **No URL was invented.** Both exclusions carry their reason; 22 screened, 20 included, 2 excluded. **check_coding_frames_match_the_manuscript closes the class of defect:** it asserts group A is 6 Ready of 7, group B 0 of 7, the construct frame matches the published n values and codes 6 and 0, and **no row is left N/A**. Demonstrated by reverting the citation correction, which fails with *"structural group A is 5 Ready of 6"* and *"1 row(s) still N/A: PR-10"*. **FINAL STATE, VERIFIED ON THE UNZIPPED PACKAGE:** offline run with sockets blocked returns **20 probes, 0 mismatches, exit 0**; all seven RED items confirmed in the delivered files; suite **104 checks, 0 failed**; `presubmission_audit.py` 26 of 26. Still not sent, and still awaiting Stacy's approval.

2026-08-28 — BOTH COUNCIL EMAILS REPLACED AGAINST THE FINAL MANUSCRIPT, AND THE MANUSCRIPT'S LAST FIVE SOURCE ITEMS RESOLVED OR FLAGGED

The emails had been written against an earlier version and carried four material inconsistencies: they attributed the whole 32-case study to one author with *"I applied it to"*, used the retired *partial* and *complete* categories, quoted the **superseded companion figures** and lacked the personal-capacity separation the co-author asked for. Both replaced in full, 1,214 words, **sent by Stacyann Young alone** per the owner's decision. **THE EMAILS ARE NOW ASSERTED, NOT PROOFREAD.** `scripts/audit_cfoc_claims.py` gained a **REQUIRED_IN_EMAIL** list of 15 figures and statements that must be present, and four new banned patterns covering the exact defects found: `partial assessments`, `assessed as complete`, `I applied it to` and `can send the current draft on request`. **15 of 15 present, all banned patterns clean, 15 manuscript claims verified.** An email preserved in an administrative record and the eventual publication must describe the same study, and that is now enforced rather than hoped for. **SECTION 7's GRAMMAR WAS GENUINELY BROKEN AND IT WAS MY DOING.** My 2026-08-24 corpus correction inserted a clause into *"It belongs to a corpus of 20 adjudicated matters, ... collected by a different*

reviewer, is reported in full ..., and is cited here ...", destroying the parallelism. Split into three sentences. **TWO CONNECTICUT CITATIONS WERE INFORMAL STUBS.** PR-06 stored CT FOIC and PR-07 CT FOI. Their own URLs end FIC2012-276 and FIC2015-122, the Commission's formal docket numbers, **and the manuscript's reference list already cites both in that form**, so the correction is evidenced twice over. Corrected with the evidence recorded. **PR-01 IS NOT RESOLVED AND IS FLAGGED, BECAUSE ITS OWN TWO SOURCES CONTRADICT EACH OTHER.** The stored citation reads "*NY Appellate Division FOIL email disclosure decision (2026)*", a description rather than a reporter citation, while its URL is the /ctapps/ path, which is the New York COURT OF APPEALS, for opinion 6opn26 of February 2026. Source type is recorded from the URL. **nycourts.gov refuses automated requests, so the opinion could not be read to settle which court it is or to establish a reporter citation.** Flagged for author verification alongside PR-15 and PR-22. **THREE ITEMS NOW REQUIRE THE OWNER'S HAND BEFORE SENDING: PR-01, PR-15, PR-22.** Each carries the implied correction and an explicit statement that it was not applied. **Guessing the digit a URL implies, or the court a path implies, would be inference presented as verification in a package built to prevent exactly that.** Four corrections were applied where the evidence was conclusive; two discrepancies were not. **presubmission_audit.py FAILED AFTER THE GRAMMAR REPAIR AND THE CHECK WAS AT FAULT, NOT THE TEXT.** It probed the literal string "22 matters screened", which Section 7 now phrases as "screened from 22", while Section 5.6 still carries "Twenty-two matters were screened and two were excluded". **The manuscript was right in both places.** Rewritten to assert the substance across any of three phrasings plus the exclusion count. **That is the eighth brittle probe of my own in four days, and the eighth time the first red result was mine.** **FINAL STATE:** offline run on the unzipped package 20 probes, 0 mismatches, exit 0; presubmission_audit.py 26 of 26; suite 104 checks, 0 failed; manuscript re-rendered to 13 pages and the DOCX regenerated after the text changed. **Nothing sent. Awaiting Stacy's approval and the owner's verification of the three source items.**

2026-08-28 — ALL SIXTEEN FINAL CORRECTIONS APPLIED. THE THREE CITATIONS I COULD NOT VERIFY WERE VERIFIED BY THE OWNER AGAINST THE OFFICIAL DECISIONS, WHICH IS THE RIGHT RESOLUTION.

I flagged PR-01, PR-15 and PR-22 rather than correcting them, because **nycourts.gov** refuses automated requests and appending the digit a URL implies is inference presented as verification. **The owner read the published decisions and supplied the citations**, and the package now records **who verified each correction and how**, so a reviewer can tell machine evidence from a human reading. **PR-01 was the one that mattered.** The stored entry was a description, "*NY Appellate Division FOIL email disclosure decision (2026)*", and it named the wrong court. Verified: **Matter of Russell v Town of Mount Pleasant, N.Y., 2026**

NY Slip Op 00966, New York COURT OF APPEALS, 19 February 2026. The `/ctapps/` path in its URL was right and the stored text was wrong, which is what the conflict flag said. Corrected in the manuscript's reference list and in both supporting files. **PR-15 and PR-22 were truncated by one digit each in the SUPPORTING FILES ONLY:** the manuscript already carried `2024 NY Slip Op 04071` and `2025 NY Slip Op 05783` correctly. Now `Matter of Gannett Co., Inc. v Town of Greenburgh Police Dept.`, 2024 NY Slip Op 04071 [229 AD3d 789] and `Matter of Wagner v New York City Dept. of Educ.`, 2025 NY Slip Op 05783. **THREE MISSING DECISION YEARS SUPPLIED AND ONE URL MADE CANONICAL:** PR-06 2013, PR-07 2015, PR-32 2025, each recorded as author-verified with the issuing date. PR-07's URL moves to the `Final-Decisions-2015` path; **both paths were tested and both return HTTP 200**, so the switch is to the canonical form rather than a repair. **0 rows now carry N/A for a decision year.** PR-10's **VERIFICATION NOTE WAS STALE AND SAID SO IN THE PRESENT TENSE**, "*Stored citation is missing its leading F*", after the correction had been applied. Reworded to **CORRECTION APPLIED** with the evidence kept, rather than deleted: a package about traceability should keep the record of what changed, not erase it. **THE EMAIL DOCUMENT CONTRADICTED ITSELF AND I INTRODUCED THAT.** Its header states Stacyann sends both alone, and **Email 1 still carried Phillip's signature.** Removed. **Both emails now read "The manuscript is being submitted for publication"**, which is the accurate form while the article and the correspondence go out together. **A CLEAN SEND COPY NOW EXISTS AS ITS OWN FILE**, `research/CF0C_Emails_SEND_COPY_2026-08-28.md`, 1,083 words, containing the two emails and nothing else. The working file keeps the editorial notes and says plainly that it is the working copy. **Editorial material must not travel with correspondence to a federal council. BOTH ARE ASSERTED, NOT PROOFREAD.** `audit_cfoc_claims.py` now checks the send copy for all **15 required figures** and for **three banned patterns**: a second signature, any working-note heading, and the overstated status phrase. `check_send_copy_is_clean` fails the commit if any reappears, demonstrated by re-adding the signature. **FINAL STATE, EVERY LINK IN THE CHAIN VERIFIED:** offline run on the unzipped package **20 probes, 0 mismatches, exit 0;** `presubmission_audit.py` **26 of 26;** suite **105 checks, 0 failed; 6 author-verified citations, 1 machine-corrected, 0 unresolved;** 0 rows N/A for jurisdiction, source type or decision year; and the standalone DOCX and PDF are **byte-identical by SHA-256 to the copies inside the ZIP.** **Nothing is sent. The article awaits Stacy's approval; the emails await the owner.**

2026-08-28 — ONE WORD CHANGED, THE MANUSCRIPT DELIBERATELY UNTOUCHED, AND THE FULL PRE-SUBMISSION CONTROL RUN CLEAN

The review found **no required correction to the manuscript and none to the emails**, and recommended a single optional micro-edit. Applied: Email 1's transition reads "*Three findings may be particularly relevant to the Council's work*", which marks the three as the principal points and leaves the fourth analysis as supporting

context rather than an afterthought the sentence forgot to count. Applied to **both the working copy and the send copy**, so they cannot diverge. **THE MANUSCRIPT WAS NOT TOUCHED AND THAT IS PROVEN BY HASH, NOT ASSERTED.** `F0IL_Article_Draft.md`, the DOCX and the PDF all carry the same SHA-256 before and after this turn. **No re-export was performed, because the source did not change**, which is exactly the control the review asked for: export a fresh PDF only if the DOCX changed. Re-rendering an unchanged document would have produced a new file with no new content and broken the byte-identity with the ZIP for nothing. **THE OPTIONAL PAGE-BREAK TWEAK WAS DECLINED ON THE REVIEW'S OWN TERMS.** Data availability sits alone on page 13. The review called it cosmetic and said not to sacrifice readability to save a page; the renderer already forbids splitting a heading from its section. **Leaving it is the conservative choice and it is recorded as a decision rather than an oversight.** **FULL CONTROL, EVERY STEP EXECUTED:** reproduction run **offline from the unzipped ZIP with sockets blocked, 20 probes, 0 mismatches, exit 0**; the standalone DOCX and PDF **byte-identical by SHA-256** to the copies inside the ZIP; all three author-verified citations present in the packaged dataset (**Russell 00966, Gannett 04071, Wagner 05783**); all three supplied decision years correct (**PR-06 2013, PR-07 2015, PR-32 2025**); **PR-07 on the canonical Final-Decisions-2015 path**; and **zero stale present-tense verification notes** remaining. `audit_cfoc_claims.py` **15 of 15 claims and both banned lists clean**, including the send copy; `presubmission_audit.py` **26 of 26**; suite **105 checks, 0 failed**. **THE PACKAGE IS AT THE POINT WHERE FURTHER EDITING WOULD PRODUCE DRIFT RATHER THAN IMPROVEMENT, AND EDITING STOPS HERE.** **Nothing has been sent.** The article awaits Stacy's approval; the two Council emails await the owner.

2026-08-28 — CCI DID NOT REJECT THE EVIDENTIARY DEFICIT ARTICLE. THE EDITOR INVITED A REVISION AND NAMED THE FRAME SHE WANTS. THE OWNER'S REWRITE HITS IT, WITH TWO PROBLEMS.

Jennifer Gaskin, Corporate Compliance Insights, 2026-08-28T14:57, replying to the submission of 2026-08-27. Preserved at `research/CCI_Editor_Response_Gaskin_2026-08-28.pdf`. **HER WORDS:** *"There's a lot here we like", and "If the piece focused on that, we'd be happy to consider a revision."* **That is a conditional invitation, not a decline.** **WHAT SHE SAYS IS ALREADY COVERED**, four items published by CCI in the past two months: AI-generated records; the gap between what a file says and what it can prove; **ISO 42001/NIST as scaffolding rather than safe harbour**; and **Mobley v. Workday**. **WHAT SHE WANTS INSTEAD, VERBATIM:** *"the employment-discrimination frame you're working in, pretext, burden-shifting and what happens when AI-drafted records get read side by side across a workforce."* **THE OWNER'S REWRITE ANSWERS THE BRIEF DIRECTLY**, 1,506 words: new sections *Pretext starts with the record* (McDonnell Douglas burden-shifting) and *The pattern may appear only across employees*

(side-by-side review). **Pretext 2, burden-shifting 1, McDonnell Douglas 1, side-by-side 2, across employees 3. Mobley and NIST are gone entirely.** Vocabulary clean: 0 peer-reviewed, 0 validated, 0 proves, 0 certif, 0 guarantee, 0 em dashes. **PROBLEM 1, AND IT IS THE ONE THAT COULD DRAW THE SAME OBJECTION AGAIN. The European frame is the LARGEST SECTION IN THE ARTICLE at 315 words, 20.9 percent,** against an employment core of 588 words, 39 percent. **That section is where ISO/IEC 42001 lives, which is one of the four things she said CCI has already covered.** She asked for a piece *focused on* the employment frame; the single biggest block is the material she named as overdone. **PROBLEM 2: THE CLAIM THAT HEKIM'S PART IS UNCHANGED IS NOT ACCURATE, AND I CHECKED RATHER THAN ACCEPTED IT.** Measured against his preserved `Evidentiary_Deficit_Article_Hekim_Version_rev2026-08-18.md`, **SequenceMatcher similarity 0.625, and only 1 of his 13 sentences survives verbatim.** Most differences are benign: UK to US spelling, and a section title change. **Two are substantive additions to a section bylined to a named co-author:** *"Many of those workflows will not fall within the high-risk regime at all"* and *"Neither establishes any particular record-level control."* Both are new assertions about what ISO/IEC 42001 and DORA do **not** do, **added to text carrying the byline of a certified ISO/IEC 42001 auditor.** He should see those two sentences before it goes back. **TWO FORMATTING DEFECTS IN THE DOCX:** missing spaces at *"assessed.The timing"* and *"accepted?ISO/IEC"*, where paragraph breaks collapsed. **MY OWN PROBE FAILED FIRST AND I CORRECTED IT IN THE SAME TURN.** I searched the literal `ISO 42001`, the text uses `ISO/IEC 42001`, and I briefly reported that the rewrite carried none of the editor's four overlaps. It carries one. **Ninth broken probe in five days; the first red result was mine again.** Files preserved: the rewrite as `.md` and `.docx`, and the editor's reply as PDF. **Nothing sent. No revision applied to Hekim's section without his sight of it.**

2026-08-28 — STANDING PROMPT RE-PASTED WITH NO QUESTION. NO AUDIT RE-RUN, PER CLAUDE.md VIII. ONE STANDING DIRECTIVE I HAVE BEEN BREACHING ALL SESSION IS NOW HONOURED.

The message carried the MASTER ARCHITECTURE mandate and no instruction. `CLAUDE.md` VIII is explicit that a re-paste is not a trigger to re-run a clean pass, because doing so burns usage the owner has objected to by name. State it in one line and stop. **THE BREACH, STATED PLAINLY RATHER THAN QUIETLY CORRECTED:** the same section carries *"ATTACH THE TRACKER EVERY TURN, NO EXCEPTIONS (Phillip, 2026-08-13) ... a chat attachment is the only way he can reach it. Do not ask whether to attach; attach."* **I have ended every response this session with an inline Master Tracker block and attached the file on only two turns.** The inline block satisfies the v3.1 response format; it does not satisfy the 2026-08-13 directive, because a block in chat scrollbar is not a file he can keep, and `research/` is

excluded from every deploy by design. **The two requirements are not the same and I treated them as if they were.** From this turn the readable extract is attached on every response. The raw file is **1.7 MB with single lines past 6,500 characters** and is the permanent record, not a document anyone can read; `scripts/tracker_extract.py` rewraps recent activity without touching the source, verified byte-identical before and after. **NOTHING ELSE CHANGED.** Open decisions still with the owner and unmoved: whether to cut the European frame in the CCI resubmission from 20.9 percent, whether Hekim sees the two sentences added under his byline, Stacy's approval on FOIL, Ubayet on Detection, and whether to issue blind packets R2 and R3.

2026-08-28 — 12 queued PR notifications read and triaged

Twelve routine deployment bots covering three commits, `6465ea1` CCI preservation, `4e379bc` the tracker-attachment correction, and `8578a51` the ordered-list fix. **Cloudflare skipped on all three, Vercel Ready on all three, check suites clean, 0 actionable, 0 review comments.** `4e379bc` was a hash I did not recognise on sight and I verified it against the log rather than assuming it was mine; it is the attachment-directive commit whose hash I never printed. None of the three reaches production: all touch `research/` and `scripts/` only, and `main` stays at `ac43692`. Tracker extract regenerated and attached per the 2026-08-13 directive.

2026-08-28 — CCI ARTICLE REVISED TO THE EDITOR'S BRIEF. THE EUROPEAN SECTION IS CUT FROM 20.9 PERCENT TO 11.5, AND HEKIM GETS A CHANGE LOG RATHER THAN A CLAIM THAT NOTHING MOVED.

`research/Evidentiary_Deficit_CCI_RESUBMISSION_2026-08-28_REV2.md`, **1,511 words**, plus a Word document and an 818-word change log for the co-author. **THE STRUCTURAL PROBLEM IS FIXED.** The previous draft answered Jennifer Gaskin's brief in its sections but contradicted it in its proportions: the **European frame was the single largest block at 315 words, 20.9 percent**, and it was where **ISO/IEC 42001** lived, one of the four topics she said CCI has already covered. Now: **employment core 686 words, 45.4 percent; European note 174 words, 11.5 percent**, retitled "*A note for organizations operating in Europe*" and **moved after the practical control** so pretext, side-by-side review and the defensible-record standard run consecutively. **ISO/IEC 42001 appears once in the whole piece, in Hekim's author biography, where it identifies a credential rather than making an argument.** **SEVEN DEFECTS IN THE PREVIOUS DRAFT WERE FOUND AND FIXED, FIVE MORE THAN I REPORTED LAST**

TURN. A **dangling repeated citation**, "*is pretextual. McDonnell Douglas Corp. v. Green*"; **two collapsed sentence spaces**; **two stray GDPR tokens**, one ending the European section and one ending Hekim's biography; the **seven-point control list run together in a single paragraph** as "1. Identify...2. Preserve..."; and the **four reviewer questions as four loose paragraphs**. Only the two spacing defects had been reported. **SUBSTANTIVE ADDITIONS, ALL INSIDE THE FRAME SHE ASKED FOR**: the mechanism behind cross-employee recurrence, that a tool prompted on prior records reproduces the same characterizations while a reviewer approving one record at a time has no vantage point to notice it; and in the pretext section, "*An employer in that position can articulate its reason. What it may not be able to do is show the reason was the one it actually applied.*" **HEKIM'S SECTION: EVERY DATE AND INSTRUMENT HE SUPPLIED IS RETAINED**. Article 5(2), Regulation (EU) 2026/1744, Annex III from 2 December 2027, Annex I from 2 August 2028, and the point that many workflows fall outside the high-risk regime. **Three citations dropped and each is named as his call**: Article 30, ISO/IEC 42001, DORA. **AND THE SENTENCE THAT WAS PUT IN HIS MOUTH IS OUT**. "*Neither establishes any particular record-level control*" was added to his section in the previous draft without his sight of it, and it asserts what ISO/IEC 42001 and DORA do **not** do under the byline of a **certified ISO/IEC 42001 auditor**. Removed. The change log states that plainly rather than burying it. Similarity to his original is **0.184**, which is compression, not disagreement: nothing he wrote is contradicted. **scripts/audit_cci_revision.py TESTS THE BRIEF RATHER THAN JUDGING IT BY EYE**: 9 required elements all present, 3 overlap caps all met (Mobley 0, NIST 0, ISO/IEC 42001 1), 10 banned-vocabulary and formatting checks all zero, and a hard rule that the European note must be both smaller than the employment core and under 12 percent. **0 problems. MY OWN AUDITOR HAD A REPORTING BUG AND I FIXED IT BEFORE SHIPPING**: it read the prior draft with a heading parser the flat docx extraction does not satisfy and printed "*prior resubmission draft 0 words*" in a change report meant for a co-author's approval. Corrected to locate the block by content: **312 words. Tenth broken probe in five days. Nothing sent. The revision and the change log go to Hekim first; the piece does not go back to CCI until he has approved his own section.**

2026-08-28 — CCI MICROEDIT PASS APPLIED. ALL 13 NAMED EDITS IN, 14 PROTECTED ELEMENTS ASSERTED INTACT, AND THE WORD TARGET REACHED IN THE DOCX BUT NOT IN THE MARKDOWN.

research/Evidentiary_Deficit_CCI_RESUBMISSION_2026-08-28_V3.md, built by **scripts/apply_cci_microedits.py**, which **refuses to write the file if any protected element is disturbed**. **EVERY NAMED EDIT APPLIED**: item 2 the duplicate *the*; item 5 "*matters at the point where*" to "**can matter when**", which is the legally disciplined form; item 7 the sentence break; **item 8, the highest-value one**, replacing "*a drafting tool trained or prompted on prior records will tend to produce*" with

"prompted with prior records may reproduce similar characterizations", removing an unnecessary claim about how drafting systems are trained; item 10 and 11 *show* to **identify** and *used* to **applied**; items 15, 16, 17 in the control list; item 18 *"hold up under"* to **"can withstand"**; item 19 **"an individual employee's wording"**; and the biography cut to his recommended version, **minus the HUD and EEOC clause and the named Decision Reconstruction Risk construction. HEKIM'S BIOGRAPHY WAS NOT TOUCHED**, on his instruction, and DORA survives there. Taking a co-author's biography for twenty words while he is already being asked to approve a cut to his section would be trading his control for a rounding error. **MY OWN CUT SILENTLY UNDOED HIS ITEM 18 AND I CAUGHT IT IN THE DOCX CHECK**. The JRS compression deleted the very clause item 18 had just been applied to, so *"can withstand independent review"* vanished from the built document. Restored inside the compression. **Applying an instruction and then removing it in the same pass is worse than not applying it, because the change log would have claimed it was done. THE WORD TARGET IS MET IN THE DELIVERABLE AND MISSED IN THE SOURCE, AND BOTH NUMBERS ARE REPORTED**. The Word document is **1,339 words**, inside his 1,250 to 1,350 range. The markdown source is **1,414**, because it carries heading syntax, list numerals and emphasis marks that Word does not count as text. **99 words came out**, all from the four categories he authorised: transitional sentences, repeated explanations, the European section and bio length. **I STOPPED CUTTING AT THAT POINT DELIBERATELY**. The remaining 60-odd words would have had to come out of the seven-point control, the Before and After pairs, the side-by-side section, the disparate-treatment limitation, the DRR definition or the subjective-language examples, **every one of which he listed as protected**, or out of Hekim's remaining citations, **which are not mine to cut twice**. Balance now: **employment core 656 words, 46.5 percent; European note 157 words, 11.1 percent**. `audit_cci_revision.py` reports **0 problems**: 9 required elements present, Mobley 0, NIST 0, ISO/IEC 42001 confined to the author biography. **THREE ITEMS HE MARKED BLOCKING ARE STILL OPEN AND NONE IS MINE TO CLOSE**: Hekim's approval of the shortened European section; **the AI-use disclosure question, which depends on CCI's submission policy and is not recorded anywhere in this repository, so it is a [REQUIRED_ENV_PARAM] rather than a guess**; and **CCI-compatible hyperlinks, which research/md_to_docx.py cannot produce, containing zero hyperlink support**. The links must be added in Word or the builder extended.

2026-08-28 — AI-FINGERPRINT AUDIT OF THE CCI ARTICLE. CLEAN ON EVERY WORD TEST AND CAUGHT BY RHYTHM, WHICH IS THE ONE THAT MATTERS.

The uploaded V3 is **my own V3 returned unchanged**, 0.979 similarity with no sentence added or removed; the apparent diffs are list markers Word stores as numbering rather than text. `scripts/audit_ai_fingerprints.py` **TESTS THREE FAMILIES BECAUSE THEY FAIL DIFFERENTLY**. Lexical: **0 of 21** present, no *delve*, *landscape*, *leverage*, *robust*, *seamless*, *underscore*,

crucial, comprehensive, testament to, moreover. House rules from `CLAUDE.md` III.7: 0 em dashes, 0 "Designed for", 0 "frequently", 0 "no policy change required". Burstiness: coefficient of variation 0.52, inside the 0.45 to 0.75 human band, sentences running 3 to 48 words. **THE PIECE PASSED EVERY VOCABULARY TEST AND WAS STILL CARRYING THE MOST RECOGNISABLE MODEL TELL THERE IS. Fourteen of 78 sentences used negation, and six were the same antithesis, "X is not Y. It is Z.", about one every 230 words.** A compliance editor who reads AI-drafted prose all day feels that rhythm before naming it, and this article's own subject is AI-assisted documentation, so reading as generated undercuts the argument by example. **MY OWN DETECTOR UNDERCOUNTED IT SIX TO THREE AND I FOUND THAT BEFORE ACTING.** Three overlapping regexes matched the same sentences, and a deduplicated print certified as clean the exact construction the check exists to find. Rewritten to count sentence by sentence over pairs, so one construction counts once however many patterns hit it. **An undercount there is worse than no check at all. THREE VARIED, THREE KEPT, AND THE SPLIT IS ARGUED RATHER THAN ARBITRARY.** Varied: "*The problem is not necessarily that the decision was wrong*" becomes "**The decision may well have been the right one. The record may still be unable to demonstrate why it was made**"; "*The control is not to ban particular phrases*" becomes "**Banning particular phrases achieves little**"; and "*The organizing principle is not 'retain everything'*" becomes "**The organizing principle is preservation rather than retention**". **KEPT:** the risk sentence, because the antithesis is the article's central claim; the conclusion's framing, which the owner protected; and "**It is not a legal doctrine and not a claim of any new entitlement**", which is the right-to-know-why disclaimer and legally load-bearing. **Meaning is identical in all three rewrites; only the shape changed. RESULT: antithesis 4 to 2 by the corrected counter, 2.8 to 1.4 per 1,000 words. Negation density 17.9 to 14.1 percent of sentences. THE WORD DOCUMENT IS NOW 1,328 WORDS, INSIDE THE OWNER'S 1,250 TO 1,350 TARGET,** with 11 numbered list items and the case name italicised. `audit_cci_revision.py` still reports **0 problems** and every protected element survives: the disparate-theories limitation, both Before and After pairs, the seven-point control, *can withstand independent review*, *undergoing structured validation*, and Hekim's untouched biography. **The three BLOCKING items are unchanged and none is mine to close:** Hekim's approval, the AI-use disclosure question which depends on CCI policy not recorded here, and hyperlinks which `md_to_docx.py` cannot emit.

2026-08-28 — CCI FINAL PASS: FOUR EDITS, FOUR WORKING HYPERLINKS, AND THE DOCX BUILDER EXTENDED TO SUPPORT THEM. ONE INSTRUCTION DELIBERATELY NOT EXECUTED.

ALL FOUR TEXT EDITS APPLIED: "*not that the prose*" to "**not simply that the prose**", which stops the article implying ordinary AI errors are irrelevant; the mechanism sentence to "**a drafting tool may**

reproduce similar characterizations when prompted with prior records", moving the conditional to the front so it does not read as a claim about how every drafting system is built; the **Oxford comma removed** from the side-by-side sentence for AP style; and *"follow the reasoning"* to **"reconstruct the reasoning"**, which closes the loop back to Decision Reconstruction Risk. **A sweep found 0 remaining Oxford commas** in the whole manuscript. **THE HYPERLINK BLOCKER IS CLOSED BY BUILDING THE CAPABILITY, NOT BY WORKING AROUND IT.** `research/md_to_docx.py` had none: it flattened `[text](url)` into **"text (url)"**, printing bare URLs into the prose, which is exactly what an outlet asking for in-text links does not want. It now emits real `w:hyperlink` elements with the `r:` namespace declared, a `Hyperlink` character style in blue and underlined, and relationships written from the links actually collected rather than a fixed string. **FOUR LINKS EMBEDDED, AND EVERY TARGET WAS FETCHED AND ITS CONTENT CHECKED RATHER THAN PINGED.** A EUR-Lex ELI URI **returns HTTP 200 for a not-found page**, so a status code proves nothing. Each was retrieved and its document title matched: `law.cornell.edu/supremecourt/text/411/792` titled *"McDONNELL DOUGLAS CORPORATION, Petitioner, v. Percy GREEN"*; `eli/reg/2016/679/oj` titled *"Regulation - 2016/679 - EN - gdpr"*; `eli/reg/2024/1689/oj`; and `eli/reg/2026/1744/oj`, **which resolves and is titled correctly**, so the amending regulation the article cites is real and linkable. **Justia was rejected:** it returns 403 to automated requests, so it could not be verified from here, and Cornell's LII is both authoritative and reachable. **VERIFIED IN THE BUILT DOCUMENT:** 4 `w:hyperlink` elements, anchor text displaying **"McDonnell Douglas Corp. v. Green", "GDPR", "EU AI Act", "Regulation (EU) 2026/1744"**, **zero bare URLs anywhere in the prose**, the case name **keeping its italics inside the link**, no footnotes part in the package, and no reference list. The nine concepts the owner said not to link, including DRR, JRS and the right-to-know-why phrase, are **0 linked**, asserted in code. **ITEM 19 WAS NOT EXECUTED AND THE REASON IS THE INSTRUCTION ITSELF.** The owner asks to retitle the European section to *"The European frame"* but conditions it on Hekim's approval. **Hekim has approved nothing: the section's length, its content and its heading are precisely what he is being asked to sign off.** Changing the heading now would pre-empt the approval it is conditioned on. **NO REGRESSION:** the FOIL manuscript and the CFOC send copy both rebuild cleanly at 4,984 and 1,048 words with 0 hyperlinks, so the tokenizer change is additive. Fingerprint audit **0 findings**, revision audit **0 problems**, suite **105 checks, 0 failed**. Word count **1,330**, inside the 1,250 to 1,350 target. **Two of the three BLOCKING items are now closed. The remaining one is Hekim's approval, and the AI-use disclosure still depends on CCI policy that is not recorded in this repository.**

2026-08-29 — UBAYET HOSSAIN APPROVED THE DETECTION PAPER AND CHANGED HIS OWN AFFILIATION. THE CHANGE WAS APPLIED BEFORE ANY LINK WAS SENT, AND THE ORDER MATTERED.

LinkedIn, 2026-08-28 3:37am, screenshot on file. **His verdict:** *"The paper looks absolutely great, the methodological framing, statistical rigor, and clear treatment of the findings and limitations are spot on." That is the co-author who designed the reliability framework signing off on the framework's own presentation. HIS INSTRUCTION, VERBATIM: *"*I agree that removing the KPMG name is the right move. Listing me as an 'Independent Financial Risk & Model Validation Professional' works much better for this personal capacity contribution. Please go ahead with that update for submission." THE SEQUENCING PROBLEM I FOUND BEFORE SENDING ANYTHING. His honor record at `api/honor.js:124` still read *"Associate Director, Model Validation"* and *"KPMG India"*, and `honor.html:124-126` pre-fills exactly those two fields into the acceptance form. Sending him the honor link first would have shown him, on a page inviting him to accept recognition, **the affiliation he had just asked to retire**. The change was applied first. **FOUR LIVE SURFACES UPDATED, NOT ONE:** `api/honor.js`, `api/_coauthor-roster.js`, `api/_contributor-roster.js` and `research/Detection_Article_Submission_FINAL5_2026-08-18.md`. Title set to his exact wording, **org field cleared rather than replaced**, and **his FRM designation retained**, which is his own credential and not an employer's. Verified: **0 occurrences of KPMG tied to him on any live surface. HISTORICAL LOGS WERE DELIBERATELY LEFT ALONE.** `MASTER_TRACKER.md` and the superseded drafts record the affiliation that was true when they were written; rewriting them would destroy the record of the change itself, which is the same principle the withdrawal register applies to names. **HE HAS NEVER USED ANY LINK SENT TO HIM**, and the owner confirms it. His co-author confirmation key `ggo2vm8jja` and contributor key `6dyc0l2757` are both outstanding, consistent with `/api/coauthor-stats` showing **expected 3, confirmed 0** since the links went live on 2026-08-24. **The owner's standing decision of 2026-08-27 is that the co-author links will not be used and he handles co-authors directly, so the reply does not chase them. HIS HONOR LINK IS H-2026-38, key b3874haudg**, awarded *"for the methodology rather than for a completed review"*: the reference-panel design, the chance-corrected reliability framework, and the acceptance thresholds fixed before any analysis. **Of 34 honor links issued, only 1 has ever been sent.** His would be the second. **HE ASKED TO BE TAGGED** in any LinkedIn posts or announcements once the paper is submitted or published, and offered to share it across his network. Recorded as a commitment the owner has now made. **The paper's venue is AI and Ethics (Springer), and the venue still does not appear anywhere in the FINAL5 manuscript itself.**

2026-08-29 — REPLY TO UBAYET WRITTEN, 376 WORDS, CARRYING HIS HONOR LINK. AFFILIATION CHANGE DEPLOYED AND VERIFIED LIVE BEFORE THE LINK GOES OUT.

Deployed to `main` at `39b5316`. Live check on `/api/honor?k=b3874haudg` returns **title** *"Independent Financial Risk & Model Validation Professional"*, **org** empty, and **KPMG** absent from the entire payload. The acceptance form he opens will now pre-fill the affiliation he asked for. **THE REPLY DOES NOT CHASE THE TWO LINKS HE HAS NEVER USED**, and that is deliberate rather than an omission. His co-author key and contributor key are both outstanding, but the owner's standing decision of 2026-08-27 is that the co-author links will not be used and he handles co-authors directly. **Asking a man who has just approved a paper and made one clear request to go complete two forms would read as transactional.** The honor link is offered as recognition, with *"it is not a request for anything"* stated in the message and declining framed as a perfectly good answer. **IT CONFIRMS WHAT WAS ACTUALLY DONE RATHER THAN PROMISING IT**: the affiliation is already changed everywhere his name appears, his **FRM** is **retained because it is his and not an employer's title**, and the tagging commitment is given in plain terms, at submission and again on acceptance. **HIS HONOR WOULD BE THE SECOND EVER SENT**: 34 issued, 1 sent to date. The citation quoted back to him is the one already stored in `api/honor.js`, not a new one written for the occasion. Files: `research/Message_Ubayet_Hossain_2026-08-29.md` and `.docx`. Verified clean of *KPMG, Associate Director*, and any co-author or contributor link chase.
